

What Checks the Proof?

A Textbook Introduction to Automated Mathematical Logic Deciders

BANK, FUTUREBANK, Imagination, Compass, Verification,
Controlled Creativity, and Formalized Self-Awareness

Textbook Version 6.80

Elementary idea. A computer can suggest useful mathematical steps, but a suggestion is not a proof. This book studies a system in which several search roles try different ways to settle a statement. Verified facts go into a shared **BANK**; imagined possibilities stay separate in a **FUTUREBANK**. A checker decides what is allowed to count as established mathematics. In short: *imagine freely; verify rigidly*.

Think of the BANK as a shared notebook containing the starting assumptions and results that have actually passed inspection. The FUTUREBANK is a scratchpad for “what if?” routes that may be worth exploring. One role tries to prove the target, another to refute it, another investigates independence, and another checks for contradiction. They may share verified discoveries, but guesses stay guesses until a certificate is checked. If no decisive certificate is found, the honest answer is UNKNOWN.

Brian Tenneson, M.A., M.S.

M.A. Mathematics, San Francisco State University

M.S. Applied Data Science, Bay Path University

btenneson2301@baypath.edu

btenneson2301.substack.com

August 22, 2026

Copyright © 2026 Brian Tenneson. All rights reserved.

What Checks the Proof? A Textbook Introduction to Automated Mathematical Logic Deciders
Textbook Version 6.80

btenneson2301@baypath.edu
btenneson2301.substack.com

Elementary Foreword

This book asks a question that does not require any formal logic to understand:

When a machine says that it has proved something, what checks the proof?

A computer can be very good at proposing ideas without every proposal being correct. That is not a defect. Human mathematicians also work by trying examples, making guesses, exploring side paths, and asking “what if?” The important distinction is between *searching for an answer* and *certifying an answer*. This book studies an architecture that tries to make that distinction explicit. Search may be flexible and imaginative; verification is deliberately strict. The recurring slogan is **Imagine freely; verify rigidly**.

The BANK and the FUTUREBANK

Imagine several careful mathematical explorers working around one shared notebook. Whenever one of them establishes a useful result and the result passes the required check, it can be written into the notebook for everybody to use. That shared notebook is the elementary picture behind the **BANK**.

The BANK begins with the assumptions that the problem has declared in advance. The book usually calls this starting set Γ . At an elementary level, Γ can simply be read as “the rules and facts we agreed to start with.” After that starting point is fixed, derived results are admitted only when the appropriate checker accepts their certificates. A useful idea does not enter merely because it is plausible, elegant, or produced by a sophisticated program.

Now give the same explorers a scratchpad. On it they may write, “Suppose this lemma were available,” “What happens if we pursue the negation?” or “This branch looks more promising than that one.” Those notes can guide later work without being treated as established mathematics. The **FUTUREBANK** plays that role. It represents possible proof futures: deductions, hypothetical lemmas, or branch histories that may be worth exploring.

The distinction is simple but central. The BANK represents what has actually survived checking. The FUTUREBANK represents what might happen. A FUTUREBANK idea can eventually lead to a real BANK deposit, but only by crossing the verification boundary.

Four ways to try to settle one statement

Ordinary automated theorem proving is often one-sided: give the machine a target and ask it to prove the target. An *Automated Mathematical Logic Decider*, abbreviated AMLD, organizes a broader settlement search around four principal roles.

The prover P tries to prove the target. The refuter R tries to prove its negation. The independence role I looks for a genuine metamathematical certificate that the target and its negation are both unprovable from the declared assumptions. The contradiction role C checks whether the admitted background itself leads to contradiction.

These are search roles, not four separate owners of mathematics. If R discovers a useful ordinary lemma while pursuing a refutation, that lemma may later help P . The shared BANK makes such reuse possible. Provenance records where a result came from, while the certificate says what the result actually establishes.

A tiny example shows the pattern. Suppose the starting information contains p and “if p , then q .” A prover proposes q by modus ponens. If the checker confirms that the inference is valid in the declared environment, q can enter the BANK as a derived result. The example is intentionally easy. It makes the architecture visible: there are admitted premises, a proposed step, a checker, and a verified deposit.

The verifier and the honest answer UNKNOWN

The verifier is intentionally less adventurous than the search machinery. It does not need to decide which idea is beautiful or which path is creative. Given the formal environment and a proposed certificate, its job is to answer a narrower question: does this certificate satisfy the declared rules for acceptance?

That narrowness is a strength. The search policy may improve, learn from failure, change priorities, or use new forms of planning without changing what counts as a valid proof. The architecture therefore separates *where to look* from *what counts once found*.

The verifier is not an oracle for all mathematical truth. Everything is relative to a declared setting: language, assumptions, inference rules, certificate formats, and checking procedures. Nor does a long failed search prove that the target is false or independent. If the machine exhausts its present resources without producing a decisive certificate, the honest computational output is UNKNOWN. A timeout describes the computation; it is not automatically a theorem about the mathematics.

Why imagination belongs here

The word *imagination* can sound strange when applied to a proof system, but the intended meaning is concrete. A chess engine can consider moves that have not been played, examine possible replies, and use those hypothetical futures to decide where to search. It does not confuse an imagined move with a move already made.

The FUTUREBANK gives proof search a similar counterfactual workspace. A **COMPASS** can rank possible directions, while a controller decides how to distribute effort. The machine can learn

from rejected attempts, notice unproductive regions, or temporarily explore what would follow from a hypothetical lemma. None of this makes the hypothetical lemma true. It only changes what the system chooses to investigate next.

Later chapters also discuss formalized self-description. The elementary idea is modest: a machine can represent facts about its own search, such as which branch it is exploring, which role has recently succeeded, or whether fallback search is receiving attention. Such information can guide control. It still does not redefine proof. “I have searched here for a long time” may be useful information about the computation; it does not imply “therefore the target is true.”

The map for the rest of the book

The technical chapters turn these pictures into definitions, state spaces, theorems, pseudocode, worked examples, exercises, and research conjectures. The notation becomes more precise, but a reader who keeps five elementary ideas in mind already has the basic map:

1. Search may try things that are not yet proved.
2. The BANK contains admitted starting material and verified results.
3. The FUTUREBANK contains represented possibilities, not automatically facts.
4. The roles P, R, I, C pursue proof, refutation, independence, and contradiction while sharing certified work.
5. The verifier is the gate; without a valid terminal certificate, the honest answer is UNKNOWN.

The proposal is therefore not that difficult mathematics becomes easy simply by giving software new component names. It is an architectural discipline: keep verified memory separate from speculative possibility, let several settlement roles reuse certified work, preserve where results came from, allow planning to guide search, and keep the acceptance standard explicit. Everything that follows develops those ideas more carefully.

Foreword

This book is about a particular question that appears simple until one tries to build a machine around it:

What checks the proof?

The question is not merely philosophical. It is architectural. A program may be brilliant at proposing lemmas, choosing search directions, imagining counterfactual proof routes, or even reasoning about its own strategy. None of those abilities, by themselves, tell us which mathematical claims are allowed to become established facts inside the machine.

The central object of the book is an *Automated Mathematical Logic Decider*, abbreviated AMLD. An AMLD is a certificate-producing architecture designed to settle a mathematical statement relative to a declared formal environment. Its four principal search roles are a prover P , a refuter R , a prover of independence I , and a prover of contradiction C . They do not merely race in isolation. They share a typed, provenance-preserving verified memory called the **BANK**. A separate **FUTUREBANK** represents possible proof futures that may be useful for planning before they are fully certified. A **COMPASS** ranks directions through that imagined future. A controller allocates attention. A verifier V is the fixed correctness boundary.

A distinction used throughout this edition is easy to state and important to preserve: *the role that discovers or submits a certificate is provenance; the semantic kind of the certificate says what it establishes*. A lemma discovered by R does not become an intrinsically “refuting” lemma, and a proof of c discovered while I is exploring a model-theoretic route is still a proof of c . Search-role tags record who found or submitted an object. Certificate semantics determine what the verifier may conclude from it.

The purpose of an AMLD is therefore broader than ordinary one-sided theorem proving. Given a target c relative to a background axiom or hypothesis set Γ , it may try to produce one of four kinds of positive evidence:

$$\Gamma \vdash c, \quad \Gamma \vdash \neg c, \quad \Gamma \not\vdash c \text{ and } \Gamma \not\vdash \neg c, \quad \Gamma \vdash \perp.$$

The third outcome requires a genuine metamathematical certificate; failure to find proofs on the first two sides is not enough. The fourth outcome says that the admitted background has failed an integrity test. If no appropriate certificate is found within the available resources, the honest result is UNKNOWN.

The architecture becomes concrete once the BANK and FUTUREBANK are treated as distinct mathematical objects. The BANK begins with Γ . It then accumulates only material that has survived the relevant verification rule. In this sense, the BANK is the machine’s verified mathematical

past. The FUTUREBANK is not a second theorem database. It is a representation of possible futures: what could happen if a particular deduction were tried, if a lemma were available, if the negation were pursued, or if attention were shifted from one role to another. This is no more “pie in the sky” than a chess engine considering hypothetical moves. The key is that a hypothetical move is not confused with a move that has actually been played, just as a hypothetical lemma is not confused with a theorem that has actually been verified.

This edition is deliberately written as a textbook rather than as a collection of research notes. The exposition proceeds deliberately. The opening chapter of each Part begins with a roadmap, and later substantial chapters use mini-roadmaps. Artificial examples are used before realistic ones. Exercises appear throughout, and solution sketches are collected before the research frontier. The conjectures are moved to the end so that the reader first learns what the architecture *is* before being asked what it *might* eventually accomplish.

There is one additional intellectual influence worth naming carefully. Max Tegmark’s 1998 paper on mathematical structures discusses “self-aware substructures” (SASs) and describes such a substructure as “something capable of some form of logical thought” [1]. Later he writes that “a substructure is self-aware if it thinks that it is self-aware” [1]. The formal self-awareness hierarchy used later in this book is *not* Tegmark’s hierarchy; Tegmark did not define Levels 0, 1, 2, ω for theorem provers. The influence is narrower: it motivates treating self-reference and formally represented self-description as legitimate structural subjects of mathematics. Here that idea is specialized to proof-search machines and separated sharply from claims about subjective consciousness.

Two earlier program documents provide historical context for the present architecture: the August 2026 automated-logical-decider research supplement and the July 2026 *Depths of a Simulation & Formalized Self-Awareness* manuscript [8, 9]. They are cited as prior program documents, not as substitutes for the independently cited logic, model-theory, automated-reasoning, security, and nonstandard-analysis literature.

The guiding sentence of the entire book is:

Imagine freely; verify rigidly.

The BANK remembers what has actually survived. The FUTUREBANK represents what might happen. The COMPASS points. The controller plans. Formal self-description may reason about that process. The declared environment admits the seed assumptions Γ ; after that seed is fixed, the verifier is the gate that decides which derived claims may enter verified mathematical memory.

Opening roadmap

Roadmap. Imagine the power of being able to tell whether a formal statement is true or false—in very informal terms. That sentence is deliberately intuitive rather than literally complete. An AMLD is not a universal truth oracle, and it cannot decide every formal statement. Relative to a declared background Γ and a fixed formal environment, its actual goal is to search for independently checkable evidence that a target c is provable, refutable, independent in the required metamathematical sense, or that the admitted background is inconsistent. If no such certificate is found under the declared conditions, the correct answer is UNKNOWN.

The power, if the architecture works well, comes from the breadth of statements that can be formalized. A target might be a theorem in pure mathematics, a software invariant, a protocol-safety property, an engineering reachability claim, a formally stated property of a scientific model, or a claim about whether an authentication system conforms to its specification. These application domains motivate AMLD research; none of them defines it.

Passwords and authentication appear later as *one worked application and one source of useful limiting examples*. They are not the central goal of AMLD, and the architecture is not being developed as a password-finding machine. Authentication is pedagogically valuable because it makes several general lessons vivid: search cannot manufacture information that is not present; a verifier may be stateful; external controls can sharply limit an otherwise powerful search procedure; and it can be more useful to search for a checkable defect in a formalized system than to search directly for a hidden secret. The same general architecture must ultimately be judged across mathematics and many other formal domains.

The book therefore proceeds from general to particular. First we define mathematical settlement and the certificate types sought by P, R, I, C . Next we build the shared BANK, the verifier boundary, state spaces, FUTUREBANK, imagination, planning, COMPASS control, and formalized self-description. Only after that foundation do we study applications—including authentication as one deliberately bounded special case—before turning to implementation, evaluation, exercises, and open conjectures.

Contents

Elementary Foreword	iii
Foreword	vi
Opening roadmap	viii
I What an AMLD Is	1
1 From theorem proving to mathematical settlement	2
1.1 The ordinary ATP problem	2
1.2 Settlement asks a larger question	2
1.3 Five honest outputs	3
1.4 An artificial four-outcome example	3
1.5 Why contradiction is operationally special	4
2 The formal environment and the forgotten starting point: Γ	6
2.1 A certificate is never valid in a vacuum	6
2.2 Γ is the seed of mathematical memory	6
2.3 Scope	7
2.4 A first walk through Γ and the BANK	7
2.5 A real formal theorem: $2 + 2 = 4$	8
3 The verifier: narrow on purpose	9
3.1 Role tags and semantic certificate kinds	9
3.2 The intrinsic verifier	10
3.3 Independence is not dual timeout	11
3.4 The verifier as a boundary, not an oracle	11

4	One AMLD cycle from start to finish	13
II	The BANK: Verified Shared Memory	14
5	Why a BANK changes the architecture	15
5.1	A race versus a community	15
5.2	Shared does not mean untyped	16
5.3	Origin views record provenance, not ownership	16
5.4	Bank state versus bank state space	17
6	Provenance and monotonicity	18
6.1	Append-only verified memory	18
6.2	Provenance is part of mathematical memory	18
6.3	The BANK as realized history	19
7	Verifier history is not the BANK	20
8	The shared lemma economy	21
8.1	Withdrawal does not consume	21
8.2	Cross-role value	21
8.3	When sharing can hurt	21
III	Imagination: From BANK to FUTUREBANK	23
9	Why chess is the right analogy	24
9.1	Chess search is counterfactual search	24
9.2	A deductive move	24
9.3	Where the analogy breaks	25
9.4	The move/commit distinction	25
10	The FUTUREBANK	26
10.1	Histories, not just endpoints	26
10.2	Tree versus DAG	26
10.3	A graded FUTUREBANK	27
10.4	FUTUREBANK does not mean truth	27

11 Planning, compass, and controlled creativity	28
11.1 First-order compass	28
11.2 Second-order compass	28
11.3 Failure as information	29
11.4 Why imagination is not mystical	29
12 Backfilling hypothetical lemmas	30
IV Reflection and Formal Self-Description	31
13 Tegmark’s SAS idea and what this book borrows	32
13.1 The Tegmark connection	32
13.2 A tagged machine	32
13.3 Why iterated provability is not used as the level scale	33
13.4 A checked self-description scale	33
14 Imagination and self-awareness are different axes	35
14.1 FUTUREBANK-reflective AMLDs	35
15 Reflection can guide search but cannot redefine proof	37
16 Safety, liveness, and fair fallback	38
16.1 Safety versus liveness	38
16.2 Fair schedules	38
16.3 Complete fallback	38
V Applications and Limits	40
17 Unlimited hypernatural budgets: what they do and do not buy	41
17.1 The nonstandard setting	41
17.1.1 Compactness supplies the unlimited horizon	41
17.1.2 Two frameworks that must not be conflated	43
17.1.3 The bounded proof-existence condition	44
17.2 Unlimited-bound proof equivalence	44
17.3 External classification, not an ordinary decider	45

17.4	Why no ordinary finite cutoff can replace an unlimited one	46
17.5	The direct-standardization barrier	47
17.6	Three conclusions to keep separate	47
18	Authentication as one worked special case	49
18.1	Why passwords belong in this book	49
18.2	A minimal authentication model	50
18.3	The right quantity is top- N probability mass	50
18.4	No-information imagination cannot create information	51
18.5	A min-entropy version of the flood-control theorem	52
18.6	Human-chosen passwords, random passwords, and password managers	52
18.7	Online and offline verification are different worlds	53
18.8	Multi-factor authentication changes the target predicate	53
18.9	The more interesting target: prove the verifier conforms to its specification	54
18.10	A four-role AMLD audit of an authentication system	55
18.11	Recovery routes and the weakest-route upper bound	55
18.12	Flood control can itself create a control problem	55
18.13	The stateful-verifier mismatch with mathematical proof	56
18.14	What an AMLD can realistically contribute to password security	56
18.15	Worked numerical example: when N makes guessing negligible	57
18.16	Exercises	57
VI	Implementation	58
19	Data structures	59
19.1	Seeding from Γ	60
20	The narrow verifier and update rule	62
20.1	Terminal classification	63
21	Imagination, compass, and selection	64
21.1	Controlled structural escape	65
22	The full AMLD loop	66
23	Safe compression and quotienting	68

24 Software invariants and tests	69
VII Worked Examples and Exercises	71
25 Four tiny worked AMLDs	72
25.1 P settles q from p and $p \rightarrow q$	72
25.2 R settles r from $\neg r$	72
25.3 I settles p over empty propositional theory	72
25.4 C detects a broken foundation	73
26 A line-by-line PA AMLD run: $1 + 1 = 5$	74
26.1 The decider and its target	74
26.2 The relevant slice of $\Gamma = \text{PA}$	75
26.3 A fair schedule chosen to make every role visible	75
26.4 Turn 1: P normalizes the left side	75
26.5 Turn 2: I considers independence	76
26.6 Turn 3: C contributes a theorem without finding inconsistency	76
26.7 The FUTUREBANK just before the winning turn	77
26.8 Turn 4: R submits the terminal refutation	77
26.9 The BANK ledger for the whole run	78
26.10 What the example teaches	78
27 A worked FUTUREBANK example	80
28 Exercises	81
29 Solution sketches	82
VIII Research Frontier	86
30 Conjectures	87
30.1 Password-domain research conjectures	88
30.2 Tempting stronger claims that should already be rejected	88
31 LAB: How to try to refute the conjectures	90
31.1 First lesson: a vague existential conjecture is hard to refute experimentally	90

31.2	Universal Lab protocol	91
31.3	Lab A: Reflective navigation advantage	91
31.4	Lab B: Graded imagination advantage	92
31.5	Lab C: Tension-triggered structural escape	92
31.6	Lab D: Backfill advantage	92
31.7	Lab E: Moderate reflection optimum	92
31.8	Lab F: Audit-over-secret advantage	93
31.9	Lab G: Cross-route shared-BANK advantage	93
31.10	Lab H: Adaptive flood-control frontier	93
31.11	Lab I: Verifier-history repair advantage	93
31.12	What counts as a successful refutation?	93
31.13	Minimal experiment pseudocode	94
	Closing roadmap	95
	References and live links	96
	Editorial provenance note	98

Part I

What an AMLD Is

Chapter 1

From theorem proving to mathematical settlement

Roadmap. We begin with the simplest distinction in the book: proving one target is not the same task as settling its status. We define four certificate-indexed settlement conclusions plus the operational output UNKNOWN, explain why timeouts are not mathematical evidence, and identify exactly what the four search roles are trying to certify.

1.1 The ordinary ATP problem

An automated theorem prover, or ATP, is usually given a formal theory and a target formula. Its success condition is simple: find a proof. If the theory is represented by Γ and the target by c , the ATP searches for a certificate of

$$\Gamma \vdash c.$$

This is a perfectly coherent task. It is also one-sided. If the ATP does not find a proof, several very different explanations remain possible: perhaps c is false relative to Γ ; perhaps c is independent of Γ ; perhaps Γ is inconsistent; perhaps a proof exists but the heuristic missed it; perhaps the proof is merely too large for the current budget.

The point is not that theorem proving is defective. The point is that *non-success has many meanings*.

1.2 Settlement asks a larger question

An AMLD is designed around a larger question:

What positive certificate, if any, settles the logical or metamathematical status of c relative to Γ ?

For this book, the principal roles are

$$\mathcal{A} = \{P, R, I, C\}.$$

Their intended terminal certificates are:

$$\begin{aligned} P &: \Gamma \vdash c, \\ R &: \Gamma \vdash \neg c, \\ I &: \text{a checked certificate that } \Gamma \not\vdash c \text{ and } \Gamma \not\vdash \neg c, \\ C &: \Gamma \vdash \perp. \end{aligned}$$

The letters name search roles, not truth values. P is a prover of c . R is a refuter of c , equivalently a prover of $\neg c$. I is a prover of an appropriate independence claim. C is a prover of contradiction.

These are *principal search objectives*, not ownership rules for all mathematics discovered during search. An agent may generate intermediate theorems, model facts, simplifications, or even a terminal certificate normally associated with another role. The BANK records the origin of such a discovery, but the verifier interprets the certificate according to its declared semantic kind.

1.3 Five honest outputs

A practical AMLD may return

PROVED, REFUTED, INDEPENDENT, INCONSISTENT, or UNKNOWN.

The first four are certificate-indexed conclusions. UNKNOWN is not a fifth mathematical truth status. It means that the machine has not produced a certificate licensing one of the other outcomes under the declared resource conditions.

Important distinction. A timeout is a fact about a computation. Independence is a fact about derivability relative to a formal setting. The former does not become the latter merely because the timeout is long.

1.4 An artificial four-outcome example

Before discussing software, it helps to see tiny instances where each role has a natural job.

Example 1.1 (A P -win). Let $\Gamma = \{p, p \rightarrow q\}$ and let $c = q$. The prover P may apply modus ponens to the two admitted premises. A verifier checks the inference and accepts q . The status is PROVED.

Example 1.2 (An R -win). Let $\Gamma = \{\neg r\}$ and let $c = r$. The refuter already has a certified premise establishing $\neg c$. The status is REFUTED.

Example 1.3 (An I -win in propositional logic). Let $\Gamma = \emptyset$ and let $c = p$. In classical propositional logic, take one valuation with p true and another with p false. The first is a model of $\Gamma \cup \{c\}$ and the second is a model of $\Gamma \cup \{\neg c\}$. By soundness, neither p nor $\neg p$ is derivable from the empty theory. A model-pair certificate can therefore establish INDEPENDENT relative to the declared propositional calculus.

Example 1.4 (A C -win). Let $\Gamma = \{s, \neg s\}$ in ordinary classical logic. The contradiction agent can certify $\Gamma \vdash \perp$. The appropriate status is **INCONSISTENT**.

These examples are intentionally small. Their purpose is to separate the meanings of the four outcomes before search complexity enters the story.

1.5 Why contradiction is operationally special

In an explosive logic, contradiction is not just another symmetric victory.

Theorem 1.5 (Contradiction dominance under explosion). *Assume the admitted logic validates explosion: from $\perp$ every formula is derivable. If C supplies a verified certificate of $\Gamma \vdash \perp$, then proofs of c and $\neg c$ cease to discriminate the intended status of c relative to that inconsistent background.*

Proof. By explosion, $\Gamma \vdash \perp$ implies $\Gamma \vdash c$ and $\Gamma \vdash \neg c$. Therefore the labels “proved” and “refuted” can both be obtained and no longer separate the two possibilities. The integrity status of the background has become the logically prior fact. $\square$

A sensible explosive-logic implementation therefore halts with **INCONSISTENT** once an accepted contradiction certificate appears. In a paraconsistent logic the rule must be reconsidered; this is one reason the logical environment must be explicit.

Corollary 1.6 (Evidence overlap under inconsistency). *Under the hypotheses of Theorem 1.5, if $\Gamma \vdash \perp$, then for every target c the three derivability conditions*

$$\Gamma \vdash c, \quad \Gamma \vdash \neg c, \quad \Gamma \vdash \perp$$

hold simultaneously. Thus the P , R , and C evidence conditions are not globally disjoint.

Proof. The contradiction condition is assumed. Explosion yields both c and $\neg c$. $\square$

This is why the words *evidence condition* and *exclusive status* should not be silently identified. An accepted proof of c is a sound certificate of derivability even if the background is later discovered to be inconsistent. To interpret **PROVED**, **REFUTED**, **INDEPENDENT**, and **INCONSISTENT** as mutually exclusive logical classes, an additional consistency condition is needed.

Theorem 1.7 (Consistent-background status partition). *Assume a classical metatheory and an object logic in which $\Gamma \vdash c$ together with $\Gamma \vdash \neg c$ entails $\Gamma \vdash \perp$. If Γ is consistent, meaning $\Gamma \not\vdash \perp$, then exactly one of the following three metatheoretic conditions holds:*

- (i) $\Gamma \vdash c$;
- (ii) $\Gamma \vdash \neg c$;
- (iii) $\Gamma \not\vdash c$ and $\Gamma \not\vdash \neg c$.

Moreover the contradiction condition $\Gamma \vdash \perp$ is false.

Proof. Classical metatheory gives the alternatives $\Gamma \vdash c$ or $\Gamma \not\vdash c$, and independently $\Gamma \vdash \neg c$ or $\Gamma \not\vdash \neg c$. Consistency rules out the case in which both derivability statements hold. The remaining possibilities are exactly (i), (ii), and (iii), and they are pairwise disjoint. By hypothesis $\Gamma \not\vdash \perp$. $\square$

Corollary 1.8 (Mutual exclusivity of sound terminal conclusions on a consistent background). *Under the hypotheses of Theorem 1.7, if the terminal checkers are sound for their declared meanings, then accepted P-, R-, and I-type terminal certificates cannot correctly certify two different members of the partition for the same fixed (Γ, c) , and a sound C-certificate cannot exist.*

Corollary 1.9 (UNKNOWN is computational, not an additional logical alternative). *Even when the hypotheses of Theorem 1.7 hold, an AMLD may return UNKNOWN under a finite resource budget. The theorem partitions the metatheoretic possibilities; it does not imply that the machine can always find a certificate for the true member of the partition.*

Exercise 1.10 (Status versus stopping). For each event, classify it as computational stopping, mathematical settlement, both, or neither: (a) a 60-second timeout; (b) an accepted proof of c ; (c) memory exhaustion; (d) a verified pair of models witnessing independence; (e) a search queue becomes empty without an exhaustive-closure certificate.

Chapter 2

The formal environment and the forgotten starting point: Γ

Mini-roadmap. This chapter introduces the object that must exist before there can be a BANK: the declared formal environment. We then make a crucial design decision: the BANK begins with explicit records for the starting axiom or hypothesis set Γ .

2.1 A certificate is never valid in a vacuum

To say that a proof is “valid” without saying relative to what is incomplete. We therefore fix a verification environment

$$\mathcal{E} = (\Gamma, c, \mathcal{L}, \mathcal{R}, \mathcal{C}_{\text{check}}, \nu),$$

where:

- Γ is the admitted axiom or hypothesis set for the instance;
- c is the target statement;
- $\mathcal{L}$ is the formal language;
- $\mathcal{R}$ is the declared inference system;
- $\mathcal{C}_{\text{check}}$ specifies certificate formats and their checkers;
- ν is a version identifier for the verification environment.

The version identifier is not decorative. If the checker, axiom base, parser, or admissibility rules change, the environment has changed.

2.2 Γ is the seed of mathematical memory

Earlier descriptions of the BANK can make it sound as though the bank begins empty and is later populated by discovered lemmas. That hides the most important first deposit. The machine has to begin from somewhere.

Definition 2.1 (Seed bank). Let $\text{seed}(\Gamma, \mathcal{E})$ be the collection of trusted initial records corresponding to the admitted axioms and hypotheses in Γ , with their scopes and environment identifier attached. The initial verified bank is

$$B_0 = \text{seed}(\Gamma, \mathcal{E}).$$

Thus the BANK starts with the formal assumptions of the problem. An axiom record is not being claimed to have been derived from earlier bank material. It is being recorded as an admitted starting premise of the environment.

Design principle. The BANK should distinguish AXIOM/HYPOTHESIS records from DERIVED records. Both may be admissible premises, but their provenance is different.

2.3 Scope

Suppose a branch temporarily assumes λ in order to investigate “what follows if λ ?” A statement proved under that assumption has scope $\Gamma \cup \{\lambda\}$. It may not be deposited as though it were proved from Γ alone.

A bank record should therefore include at least:

claim; certificate or admission tag; semantic kind; origin; scope; dependencies; provenance; verifier record; deposit time; environment ID.

This small amount of bureaucracy prevents large logical errors later.

Here *origin* and *semantic kind* are deliberately separate fields. Origin answers a historical question: “which search role submitted this record?” or “was this premise admitted by the environment?”

Semantic kind answers a logical question: is this an admitted hypothesis, an intermediate theorem, a target proof, a negated-target proof, an independence certificate, or a contradiction certificate?

2.4 A first walk through Γ and the BANK

Return to $\Gamma = \{p, p \rightarrow q\}$ and $c = q$. The machine begins with

$$B_0 = \{[p : \text{HYP}], [p \rightarrow q : \text{HYP}]\}.$$

The prover reads both records, proposes the inference “modus ponens yields q ,” and submits its certificate. If V accepts it, the next state is

$$B_1 = B_0 \cup \{[q : \text{DERIVED}]\}.$$

The BANK did not merely store the answer; its dependency and provenance fields preserve the route by which the answer became admissible.

2.5 A real formal theorem: $2 + 2 = 4$

The Metamath Proof Explorer contains the formally checked theorem `2p2e4`, whose assertion is

$$(2 + 2) = 4.$$

Its live theorem page displays the formal proof and dependencies [3]. This provides a concrete example of what an AMLD terminal event can look like.

Imagine an AMLD environment whose language and foundational assumptions are those of the relevant `set.mm` context, whose target is $c \equiv (2 + 2 = 4)$, and whose verifier is a compatible Metamath proof checker. The P role submits the existing proof certificate. The verifier checks it. If accepted, the AMLD returns `PROVED`.

This is a *real formally verified mathematical statement*, not merely the propositional toy example. At the same time, one should be precise about what is and is not being claimed: this paragraph describes a valid AMLD run at the architectural level using an existing Metamath certificate. It is not a claim that a particular current AMLD implementation independently rediscovered the published proof.

The example is pedagogically useful because it makes the boundary visible. Search may be trivial if the certificate is handed to P ; verification is still meaningful. In harder instances the search problem becomes the dominant cost, but the acceptance semantics need not change.

Exercise 2.2 (Seeding the bank). Let $\Gamma = \{a, a \rightarrow b, b \rightarrow d\}$ and target d . Write the initial bank records and a two-deposit history (three BANK states B_0, B_1, B_2) that ends with d . Which records are admitted premises and which are derived?

Chapter 3

The verifier: narrow on purpose

Mini-roadmap. We now isolate the verifier from the rest of the machine. The verifier should know enough to check a certificate and its admissible dependencies, but it should not care whether the compass likes the certificate or whether a reflective controller thinks the search is going well.

3.1 Role tags and semantic certificate kinds

Two classifications are needed, and they should not be conflated.

The *submitting role* is one of

$$\mathcal{A} = \{P, R, I, C\}.$$

Separately, let $\mathcal{T}$ be the set of declared semantic certificate kinds. It may include terminal kinds such as

TARGET_PROOF, NEGATED_TARGET_PROOF, INDEPENDENCE, CONTRADICTION, as well as nonterminal kinds such as INTERMEDIATE_THEOREM, MODEL_FACT, or other environment-specific certificates.

For each semantic kind $t \in \mathcal{T}$, let $\mathcal{K}_t$ be its certificate-body space and define the semantically tagged certificate space

$$\mathcal{K}_{\text{sem}} = \bigsqcup_{t \in \mathcal{T}} (\{t\} \times \mathcal{K}_t).$$

A submitted proposal then lies in the role-tagged proposal space

$$\mathcal{K} = \bigsqcup_{a \in \mathcal{A}} (\{a\} \times \mathcal{K}_{\text{sem}}).$$

Thus a proposal has the conceptual form

$$k = (a, (t, \kappa)),$$

where a says *who submitted it*, t says *what kind of certificate it claims to be*, and κ is the certificate body. For this abstract interface, κ is understood to contain or bind every truth-relevant field needed by the declared checker—in particular the claimed conclusion, scope, dependency references, and proof/metacertificate data as appropriate to kind t . Implementations may store those fields separately while passing the same semantic content to the checker.

Theorem 3.1 (Unique submitting-role tag). *Every proposal $k \in \mathcal{K}$ has a unique submitting role $a \in \{P, R, I, C\}$ and a unique semantically tagged component $s \in \mathcal{K}_{\text{sem}}$ such that $k = (a, s)$.*

Proof. This is exactly the defining property of the tagged disjoint union. Every element belongs to one tagged component, and different role-tagged components are disjoint. $\square$

Corollary 3.2 (No untyped agent proposal). *If the AMLD selector is required to output elements of $\mathcal{K}$, every submitted proposal identifies exactly one of the four search roles. There is no fifth untagged agent proposal unless the architecture explicitly enlarges $\mathcal{A}$.*

The theorem is syntactic and architectural. It does *not* say that every theorem in the BANK is intrinsically a P -theorem, R -theorem, I -theorem, or C -theorem. A commutativity lemma found by R remains a commutativity lemma. Its origin tag records history; its semantic kind records what the verifier established.

Let

$$\tau : \mathcal{T}_{\text{terminal}} \longrightarrow \{P, R, I, C\}$$

map each terminal semantic kind to its designated settlement role: target proof to P , negated-target proof to R , independence to I , and contradiction to C . The designated settlement role need not equal the submitting role.

Corollary 3.3 (Terminal-role coverage). *If the terminal semantic kinds are exactly the four kinds in the domain of τ , every accepted terminal certificate has a designated settlement role in $\{P, R, I, C\}$.*

Remark 3.4 (Submitting role versus settlement role). Suppose R , while exploring a refutation, unexpectedly constructs a valid proof of c . The proposal may retain submitting-role tag R for provenance while carrying semantic kind `TARGET_PROOF`. If the checker accepts the certificate, the semantic settlement role is P and the derivability conclusion is `PROVED`. Nothing mathematically requires the discovery to be thrown away or relabeled as having been found by P .

Exercise 3.5 (Origin versus semantic kind). An I -agent, while trying to build two models, discovers an ordinary object-level lemma L and later a valid proof of the target c . Give appropriate submitting-role tags and semantic kinds for both proposals. If both are accepted, what belongs in provenance, what belongs in theorem memory, and what terminal status does the second certificate license?

3.2 The intrinsic verifier

Let $\Gamma_{\mathcal{E}}(\mathbf{B})$ mean the base environment together with the verified bank records admissible in the certificate's scope. The intrinsic checker acts on semantic certificate content, not on the submitting-role tag. Define

$$v_{\mathcal{E}} : \mathfrak{B} \times \mathcal{K}_{\text{sem}} \longrightarrow \{0, 1\}.$$

A software wrapper may receive the full proposal $k = (a, s)$, record a as provenance, resolve the admissible dependencies named by s , and pass the truth-relevant semantic content to the intrinsic checker. In a concrete implementation the fields making up s may be stored separately—for example as semantic kind, claim, scope, dependencies, and certificate body—but the submitting-role tag remains outside the intrinsic truth predicate. The value 1 means that the submitted certificate is accepted. Real software may return richer diagnostics such as `REJECT`, `MALFORMED`, or `ERROR`; only acceptance authorizes a mathematical deposit.

Theorem 3.6 (Policy-neutral verification). *Fix $\mathcal{E}$, a verified bank state $\mathbf{B}$, and a semantic certificate $s \in \mathcal{K}_{\text{sem}}$. If the intrinsic verifier is defined only from $(\mathcal{E}, \mathbf{B}, s)$, then changing the *FUTUREBANK*, search history, compass, controller, novelty score, formal self-model, or submitting-role tag cannot change the truth value of “ s is accepted in $\mathbf{B}$.”*

Proof. Those coordinates are not arguments of the acceptance predicate. They may change which semantic certificate is selected, or who submits it, but not the verifier’s value on a fixed input triple. $\square$

This theorem is almost tautological. That is a feature. Good architecture often turns a difficult safety requirement into a simple type or interface fact.

Corollary 3.7 (Submitting-role neutrality). *Let $a, b \in \mathcal{A}$ and let $s \in \mathcal{K}_{\text{sem}}$. If two proposals differ only in submitting-role tag, (a, s) versus (b, s) , then their intrinsic verification result is identical in the same environment and admissible bank state.*

Proof. The intrinsic verifier receives the same semantic input s in both cases. The role tag is recorded as provenance but is not an argument of $v_{\mathcal{E}}$. $\square$

Corollary 3.8 (Discovery-role irrelevance for terminal semantics). *Under the same interface, if any search role submits a semantic certificate of kind *TARGET_PROOF* and the verifier accepts it, the mathematical conclusion is $\Gamma \vdash c$; analogously for the other terminal semantic kinds. Who discovered the certificate does not change what it proves.*

3.3 Independence is not dual timeout

The *I* role is the least like the other three. A proof of c , a proof of $\neg c$, and a proof of $\perp$ can all be object-level derivations. A proof of independence usually lives one level up.

In a model-theoretic setting, one sufficient pattern is a pair

$$M^+ \models \Gamma \cup \{c\}, \quad M^- \models \Gamma \cup \{\neg c\}.$$

Under the appropriate soundness theorem, the first model rules out $\Gamma \vdash \neg c$ and the second rules out $\Gamma \vdash c$. In other settings, *I* may provide a formalized relative-consistency argument, a metatheoretic derivation, or an exhaustive finite certificate. The certificate type must be declared in $\mathcal{C}_{\text{check}}$.

The TPTP/SZS ecosystem contains a mature vocabulary for automated-reasoning outcomes, distinguishing statuses such as theoremhood, counter-satisfiability, satisfiability, and unsatisfiability, and it supports proof/model output conventions [4, 5]. The AMLD model-pair pattern should therefore be understood as building on established proof- and model-producing reasoning practice, not as if model-based negative evidence had no precedent.

3.4 The verifier as a boundary, not an oracle

A useful verifier is strict and narrow. It does not answer “is this statement really true in all possible senses?” It answers a declared formal question: does this certificate satisfy the checker for this environment and these dependencies?

That is enough to support a strong internal epistemology:

The machine may be uncertain about what to try next while being exact about what it has already certified.

Chapter 4

One AMLD cycle from start to finish

Mini-roadmap. The pieces are now sufficient to describe a complete cycle: seed the BANK from Γ , choose a role, read verified material, propose an intermediate or terminal certificate, verify it, update state, and either halt or continue.

At step m , let the full state eventually have the form

$$\Sigma_m = (B_m, F_m, H_m, \text{Ref}_m, \text{Ctrl}_m).$$

Let $\mathfrak{S}$ denote the declared set of admissible tuples of this form—that is, tuples satisfying the architecture’s typing, scope, environment, and state-invariant requirements—so $\Sigma_m \in \mathfrak{S}$. The individual coordinate spaces may vary by implementation; the notation $\mathfrak{S}$ does not assume that every arbitrary tuple is admissible.

For the moment ignore all but B_m and H_m .

One turn is:

1. choose an active search role $a \in \{P, R, I, C\}$;
2. allow a to read any admissible verified records in B_m ;
3. let a propose one new certificate or certified intermediate result;
4. submit it to V ;
5. if accepted, add the resulting typed, scoped record to the BANK;
6. always record the attempt and diagnostics in verifier/search history;
7. if the accepted object is terminal, report the corresponding status and halt;
8. otherwise continue.

The distinction between *intermediate* and *terminal* accepted records is important. The BANK can grow for a long time without the target being settled.

Checkpoint. At this point the reader should be able to explain an AMLD without using the words FUTUREBANK, compass, creativity, or self-awareness. Those ideas improve search. They are not required to define the basic correctness boundary.

Part II

The BANK: Verified Shared Memory

Chapter 5

Why a BANK changes the architecture

Roadmap. A bare conjecture settler can run several searches and wait for one to return a decisive certificate. The BANK makes the architecture cooperative and cumulative. This Part explains what is genuinely gained, what is not new in the broader literature, and which invariants make shared memory safe.

5.1 A race versus a community

Imagine four isolated processes. P searches for c , R for $\neg c$, I for independence, and C for contradiction. When any one succeeds, the system halts. That is a multiway conjecture settler, but the searches need not teach one another anything.

Now add a verified shared BANK. A lemma found while R is pursuing $\neg c$ may later help P . A model-theoretic construction found by I may expose a useful object-level lemma. A simplification found while C is testing consistency may help every role. The history of verified work becomes reusable.

The difference is not merely parallelism. It is *cross-role accumulation*.

Bare multiway settler	BANK-centered AMLD
Roles may search independently.	Roles read a common verified memory.
Intermediate work may disappear with a branch.	Accepted lemmas persist with provenance.
A role chiefly benefits itself.	A discovery by one role may help another.
The state is mainly “who is searching where?”	The state includes an accumulating mathematical object.
Search history and theorem memory may blur.	Verified BANK and verifier history are separated.

This use of shared memory has important precedents in blackboard architectures and cooperative theorem proving. For example, work on Leo-III explicitly explores an agent-based blackboard architecture [6]. The claim made here is therefore not that shared blackboards were invented

by AMLD. The architectural novelty *within this program* is the way a certificate-gated, typed, provenance-aware BANK is made central to four-way mathematical settlement and later paired with a quarantined FUTUREBANK.

5.2 Shared does not mean untyped

The BANK is shared, but a metamathematical independence statement is not automatically usable as an object-level theorem inside the theory whose independence is being discussed. A scoped hypothetical lemma is not automatically global. A contradiction certificate changes the interpretation of the whole run in explosive logics.

Hence the BANK is both shared and typed.

5.3 Origin views record provenance, not ownership

The BANK itself is best treated as one set of verified records. Agent-specific “compartments” are useful indexes or views, but they should not define the ontology of the BANK because the initial Γ records were not discovered by any of P, R, I, C .

Let the origin-label set be

$$\mathcal{O} = \{\text{ENV}, \text{IMPORT}, P, R, I, C\}.$$

Here ENV marks axioms or hypotheses admitted directly by the current verification environment, while IMPORT may mark separately imported, already certified material. For $o \in \mathcal{O}$, define the origin view

$$B_{o,m} = \{r \in B_m : \text{origin}(r) = o\}.$$

These views say where records came from. They do not say who owns them, and they do not determine their semantic kind.

Proposition 5.1 (Origin-view partition). *If every BANK record has exactly one origin label in $\mathcal{O}$, then the origin views are pairwise disjoint and*

$$B_m = \bigsqcup_{o \in \mathcal{O}} B_{o,m}.$$

Proof. Every record has one origin, so it belongs to one origin view. Uniqueness of origin prevents membership in two distinct views. $\square$

Corollary 5.2 (Seed records need no fictitious agent). *Every initial record produced by $\text{seed}(\Gamma, \mathcal{E})$ may have origin ENV. Thus the theorem that every agent proposal has one of the tags P, R, I, C does not force the axioms of Γ to be attributed to an agent.*

Important distinction. There are now two different classifications available for indexing. One may group records by *origin* (ENV, IMPORT, P, R, I, C), or by *semantic kind* (hypothesis, intermediate theorem, target proof, independence certificate, and so on). These classifications cross-cut each other. The same semantic kind may be discovered by different agents, and one agent may discover many semantic kinds.

5.4 Bank state versus bank state space

A state is one possible current bank. A state space is the collection of all bank states the machine could occupy. Let $\mathcal{R}_{\text{bank}}$ be the universe of well-formed BANK records for the fixed environment. A natural abstract state space is

$$\mathfrak{B} \subseteq \mathcal{P}_{\text{fin}}(\mathcal{R}_{\text{bank}}), \quad \mathbf{B}_m \in \mathfrak{B},$$

where the subset restriction enforces the environment's typing, scope, provenance, and verification invariants. Origin views and semantic-kind indexes can then be derived from the record set rather than built into the mathematical identity of the BANK.

This notation is worth keeping clean because the full AMLD later has several nested state spaces.

Chapter 6

Provenance and monotonicity

Mini-roadmap. This chapter treats the BANK as a growing proof object. We show that accepted memory is monotone under the ordinary append-only design, and explain why provenance should survive even when two branches reach the same formula.

6.1 Append-only verified memory

The simplest BANK never deletes established records during a run. Search indices may be pruned, caches may be evicted, and representations may be compressed, but the logical content is monotonically accumulated.

Theorem 6.1 (Verified-bank monotonicity). *Suppose the update rule only adds verifier-accepted records and never removes the logical availability of earlier accepted records. Then*

$$B_m \subseteq B_{m+1}$$

in the natural record-set interpretation for every m .

Proof. By the update rule, B_{m+1} is either B_m itself or B_m together with one or more newly accepted records. No old record is removed. $\square$

This theorem is simple, but it has practical consequences. It allows dependency graphs, incremental indexing, replay, and audit trails to be defined against a stable growing object.

6.2 Provenance is part of mathematical memory

For every derived record r , store a dependency set $\text{dep}(r)$. The dependency DAG can answer questions such as:

- Which axioms of Γ does this lemma ultimately use?
- Which role first produced the certificate?
- Which earlier bank records are needed to replay it?

- Does a supposedly global theorem actually depend on a temporary assumption?
- If a checker version is retired, which records must be revalidated?

The BANK can therefore be understood as a proof-carrying shared memory rather than a bag of formulas. In the record-set model used here, separately accepted derivations of the same displayed formula may remain distinct records with distinct receipts and origins; a formula-level retrieval index may still present a canonicalized view without erasing that ledger provenance.

6.3 The BANK as realized history

Write $B(0), B(1), \dots$ for snapshots through time. The sequence is itself a history. A ledger

$$L_B \subseteq \text{RecordID} \times \mathbb{N} \times \mathcal{P}_{\text{fin}}(\text{RecordID})$$

can record deposit times and dependencies.

This distinction will matter when we meet the FUTUREBANK. The BANK preserves the realized, verified past. The FUTUREBANK preserves represented possible futures.

Exercise 6.2 (Provenance error). A branch assumes h temporarily, derives u from $\Gamma \cup \{h\}$, and then deposits u as though its scope were only Γ . Explain exactly what is wrong. Give one way a bank schema could make this error mechanically detectable.

Chapter 7

Verifier history is not the BANK

Mini-roadmap. Rejected candidates can be useful data. The danger is allowing “useful” to drift into “proved.” We therefore maintain a separate history H that search may learn from but deduction may not treat as theorem memory.

Let H_m contain events such as:

- accepted and rejected proposals;
- parser and type failures;
- duplicate candidates;
- timing and resource data;
- which FUTUREBANK branch generated a proposal;
- compass scores at selection time;
- fallback invocations;
- controller actions.

A rejected lemma can lower the priority of similar candidates or trigger a structural escape. It cannot be used as a premise merely because it was interesting.

Theorem 7.1 (Verifier-gated noncontamination). *Assume (i) B_0 contains only admitted seed records and validly certified imported records, (ii) every derived entry into B requires acceptance by a sound verifier for its declared environment, and (iii) no record in H or the FUTUREBANK can bypass that gate. Then every finite bank state contains only admitted or validly certified records of the declared types and scopes.*

Proof. Induct on update steps. The base case is assumption (i). At the next step, rejected or merely speculative material changes no verified record. If a new derived record is added, assumption (ii) supplies its certification. Thus the invariant is preserved. $\square$

Design principle. Learning from failure is allowed. Reclassifying failure as proof is not.

Chapter 8

The shared lemma economy

Mini-roadmap. We now look at the BANK dynamically: agents withdraw by reading, deposit by verification, and sometimes produce lemmas whose value is discovered by a different role than the one that generated them.

8.1 Withdrawal does not consume

If P uses a bank theorem, the theorem remains available to R, I, C . The BANK is not a resource pile from which facts are removed. It is closer to a shared library with typed borrowing rights.

8.2 Cross-role value

Suppose R attempts to prove $\neg c$ and proves a general lemma L along the way. The verifier accepts L . Later P discovers that L is useful in a completely different derivation of c . This is not paradoxical. Intermediate mathematics need not inherit the final stance of the role that discovered it.

This is one reason agent names should be treated as provenance tags rather than epistemic tribes.

Corollary 8.1 (Cross-role reuse under origin-neutral admissibility). *Suppose a verified BANK record r is admissible for a later inference exactly when its semantic kind, scope, dependencies, and environment satisfy the declared use rules, and suppose those rules do not restrict use merely because of $\text{origin}(r)$. Then every search role may use r whenever those semantic admissibility conditions hold, regardless of which role originally submitted it.*

Proof. By hypothesis, admissibility is a function of the stated semantic conditions and not of origin. Changing only the role that originally discovered r therefore cannot change whether the record is admissible. $\square$

8.3 When sharing can hurt

Shared memory can increase retrieval costs, create distraction, or flood the selector with irrelevant but valid lemmas. Correctness and efficiency must therefore be separated. A BANK may be logically

safe but computationally badly indexed.

Machine-learned relevance guidance can help rank large theorem libraries, and systems such as ENIGMA illustrate how learned ranking can guide symbolic theorem proving without making the learned model itself the final proof checker [7]. This is closely aligned with the AMLD principle that guidance may change search while verification remains independently defined.

Exercise 8.2 (Cross-role reuse). Construct a toy example in which a lemma first derived by R is later useful to P . The lemma need not help R reach $\neg c$. Identify its scope and explain why agent origin does not restrict reuse.

Part III

Imagination: From BANK to FUTUREBANK

Chapter 9

Why chess is the right analogy

Roadmap. The word “imagination” can sound mystical when attached to a theorem prover. Chess removes the mystery. A chess engine considers moves that have not been played, follows their consequences, evaluates branches, and then chooses an action. We translate that idea into deductive search.

9.1 Chess search is counterfactual search

A chess engine at position x considers a legal move a , computes the successor $T(x, a)$, then considers replies from that position, and continues recursively. Claude Shannon’s classic 1950 paper helped establish chess as a model problem for symbolic machine reasoning [2].

The engine’s internal branch

$$x \rightarrow x_1 \rightarrow x_2 \rightarrow \cdots$$

does not mean those moves have occurred on the physical board. It means they are represented possibilities used for planning.

9.2 A deductive move

For theorem proving, a proof-search state x can include goals, local assumptions, substitutions, available bank records, and unresolved obligations. A *deductive move* is an action that changes this proof-search state by applying or hypothesizing a mathematically meaningful step.

Examples include:

- apply a rule of inference to selected premises;
- rewrite a goal using a bank theorem;
- introduce a quantified witness candidate;
- split into cases;
- introduce a temporary assumption for conditional reasoning;

- propose an auxiliary lemma;
- switch from object-level proving to model search;
- backtrack or change the active role.

The analogy is strong: a chess move changes a position; a deductive move changes a proof state. A sequence of moves creates a possible game; a sequence of deductive moves creates a possible proof history.

9.3 Where the analogy breaks

A chess engine can cheaply check whether many moves are legal. Mathematical proof search often contains a more expensive distinction. The system may want to explore

“If lemma λ were available, would the target become easy?”

before investing the full cost required to prove λ . Thus some proof-search branches are conditional plans rather than fully certified deduction sequences.

This motivates a graded FUTUREBANK.

9.4 The move/commit distinction

It is helpful to separate three verbs:

1. **Imagine**: represent a possible deductive move or hypothetical lemma.
2. **Attempt**: turn some imagined object into an actual certificate proposal.
3. **Commit**: after verifier acceptance, deposit the resulting record in the BANK.

A chess engine can imagine $Qh7+$ without playing it. An AMLD can imagine “prove λ , then rewrite by λ ” without claiming λ is a theorem.

Exercise 9.1 (Chess-to-proof translation). Give a chess-search concept and a proof-search analogue for each of the following: position, legal move, principal variation, transposition, evaluation function, and move actually played. Explain one place where the analogy is imperfect.

Chapter 10

The FUTUREBANK

Mini-roadmap. The FUTUREBANK is the provenance-preserving representation of possible proof futures. We define it first as a set of histories, then explain why a DAG is often a better implementation, and finally separate speculative from locally checked futures.

10.1 Histories, not just endpoints

Let $\text{Succ}(x)$ be the represented successors of proof state x . If h is a history, let $\text{end}(h)$ be its final state. Starting from the current history h_0 , define the generation layers

$$\mathcal{G}_0 = \{h_0\},$$

and recursively let $\mathcal{G}_{d+1}$ contain one-step extensions of histories in $\mathcal{G}_d$. The depth- d FUTUREBANK is

$$\mathcal{F}_d = \bigcup_{j=0}^d \mathcal{G}_j.$$

The emphasis on histories matters because two routes can reach the same visible endpoint while carrying different assumptions, dependency histories, or future options.

10.2 Tree versus DAG

A literal history tree preserves provenance automatically but can duplicate enormous downstream computation. A DAG can share equivalent future states, but only if the equivalence relation preserves everything relevant to future search and verification.

For clarity, the implementation pseudocode later uses single-parent history nodes as its correctness-first baseline. Computation sharing is then an optimization layered on top of those histories, not a reason to erase them. A FUTUREBANK implementation may share downstream work while separately retaining whatever path multiplicity, incoming transition data, or ancestry the task needs.

10.3 A graded FUTUREBANK

Distinguish

$$F_{\text{spec}}, \quad F_{\text{loc}}, \quad B.$$

Interpret them as follows:

- F_{spec} : explicitly speculative branches; useful for planning, no theorem status;
- F_{loc} : branches that have passed a declared local admissibility or transition checker, but are not necessarily complete theorem certificates;
- B : material admitted or accepted by the final verifier.

One common trust path is

$$F_{\text{spec}} \longrightarrow F_{\text{loc}} \longrightarrow V \longrightarrow B,$$

where the local-check stage is optional: a speculative object may also be submitted directly to the final verifier when its certificate format permits. The arrows indicate increasing certification, not increasing subjective confidence.

10.4 FUTUREBANK does not mean truth

This is the defining quarantine rule. A node can have enormous compass score and still be speculative. A branch can be internally coherent and still depend on an unproved lemma. A controller can be nearly certain that a route will work and still be wrong.

Theorem 10.1 (FUTUREBANK noncontamination). *Under the hypotheses of Theorem 7.1, arbitrary additions to F_{spec} cannot by themselves change the set of verified mathematical records in B .*

Proof. Speculative membership is not a deposit rule. Only verifier acceptance changes verified mathematical memory. $\square$

Chapter 11

Planning, compass, and controlled creativity

Mini-roadmap. Imagination answers “what could happen?” Planning asks “what should we try?” The compass supplies ranking information, while a controller allocates finite attention. Creativity appears as controlled structural departure, not as permission to relax proof standards.

11.1 First-order compass

A first-order compass can score mathematical transitions:

$$K_1(x, y) \in \mathbb{R}.$$

High score means that moving from proof state x toward y appears promising under the learned or hand-designed navigation model.

The score may use features such as goal similarity, prior proof trajectories, quotient-state density, dependency structure, or resource estimates. None of these are truth predicates.

11.2 Second-order compass

Once the search itself is adaptive, a higher-level score can rank search actions:

$$K_2(\Sigma, a).$$

Actions might include:

- continue the current local branch;
- expand a different FUTUREBANK region;
- switch from P to R or I ;
- increase or decrease novelty pressure;

- backfill a speculative lemma;
- compress a quotient state;
- invoke complete fallback.

Thus K_1 points through mathematics while K_2 points through search policy.

11.3 Failure as information

Suppose K_1 repeatedly favors one region, but certificates generated there repeatedly fail verification. The disagreement between guidance and observed verification history is evidence about search quality. It can justify leaving the current neighborhood.

Controlled creativity can therefore be understood as a policy that increases willingness to explore structurally different futures when some combination of stagnation, novelty opportunity, and compass/failure tension crosses a threshold.

Design principle. Creativity belongs in IMAGINE, COMPASS, SELECT, and controller state. It does not belong inside the verifier’s acceptance rule.

11.4 Why imagination is not mystical

The phrase “the machine imagines a proof” should be read operationally:

It recursively constructs and evaluates counterfactual deductive histories without confusing those histories with verified mathematics.

That is close in spirit to ordinary game-tree search. The mathematical challenge is not to make imagination metaphysical. It is to make counterfactual search useful, typed, provenance-preserving, and safe.

Chapter 12

Backfilling hypothetical lemmas

Mini-roadmap. A speculative branch can be valuable because it tells the machine what would be worth proving. We now show how a conditional plan can be converted into a verified route by proving its unmet dependencies in order.

Suppose a branch has the form

if λ_1, λ_2 are available, then c becomes reachable.

The FUTUREBANK may represent this branch even when λ_1 and λ_2 are not in the BANK. The controller can then turn them into subgoals.

Proposition 12.1 (Dependency-closed promotion). *Let a speculative branch depend on a finite acyclic set of hypothetical lemmas D . Suppose there exists an order $\lambda_1, \dots, \lambda_n$ of D such that each λ_j has a verifier-acceptable certificate from Γ together with earlier promoted dependencies $\lambda_1, \dots, \lambda_{j-1}$. Then the hypothetical dependencies can be promoted one by one to the BANK, after which the branch no longer depends on those hypotheses.*

Proof. Induct on the dependency order. The first certificate is valid from the admitted bank. After it is accepted and deposited, the second's prerequisites are available, and so on. Acyclicity ensures that the process has a topological order. $\square$

This proposition does not say that every promising imagined branch can be backfilled, nor does it claim that its stated dependency order is the only possible route. It gives a clean sufficient condition under which a speculative finite dependency DAG can be promoted into verified mathematics.

Exercise 12.2 (Backfill). A speculative route to c requires lemmas λ and μ , where μ can be proved only after λ . Draw the dependency DAG. Describe the difference between (i) exploring consequences of both lemmas speculatively and (ii) promoting them into the BANK.

Part IV

Reflection and Formal Self-Description

Chapter 13

Tegmark's SAS idea and what this book borrows

Roadmap. This Part adds formal self-description without confusing it with imagination. We begin with Tegmark's notion of self-aware substructures as an intellectual influence, then define a checked architectural scale for AMLD self-description and explain why iterated provability is not an adequate level counter. Finally we show that self-description may steer search while leaving proof semantics unchanged.

13.1 The Tegmark connection

In *Is “the theory of everything” merely the ultimate ensemble theory?*, Tegmark discusses mathematical structures complex enough to contain “self-aware substructures” and explicitly associates them with logical thought [1]. His project is cosmological and philosophical, not a design manual for automated theorem proving.

The relevance here is structural. A formal machine may contain a representation of itself and may reason about that representation. Once this possibility is admitted, one can ask precise questions: Which self-descriptions can it check? Which formal provability claims can it establish in a declared theory? Which facts about its own current search state can it safely use as controller input?

Important distinction. The Level-0, Level-1, Level-2, and Level- ω hierarchy below is a Tenneson-style formalization inspired by self-referential themes; it is not a hierarchy proposed by Tegmark. Nor is any level here a theorem about subjective consciousness.

13.2 A tagged machine

Let A be an AMLD with a formal name or code $\langle A \rangle$ available in a declared metatheory T . A minimal static self-description may contain a formula such as

$$\text{AMLD}(\langle A \rangle),$$

together with machine-checkable facts identifying the architecture version, verifier version, and role set. The point is not that the string is intrinsically profound. The point is that the formal environment can refer to the machine itself as an object of reasoning and can distinguish a checked self-description from an unchecked assertion about itself.

13.3 Why iterated provability is not used as the level scale

Suppose the declared metatheory T is recursively axiomatized strongly enough to arithmetize syntax and has a standard provability predicate Prov_T . A tempting definition would set $R_0(A) = \text{AML}(\langle A \rangle)$ and then iterate

$$R_{n+1}(A) = \text{Prov}_T(\ulcorner R_n(A) \urcorner).$$

That notation is useful in provability logic, but it is a poor measure of increasing architectural self-description. For standard arithmetized provability predicates satisfying the usual derivability conditions, if T proves a sentence A , then T also proves the corresponding assertion $\text{Prov}_T(\ulcorner A \urcorner)$. Repeating the operator can therefore add syntactic layers without adding a genuinely new self-model. Reflection principles also interact with Gödelian and Löbian limitations and require the object theory, metatheory, coding, and provability predicate to be specified with care [18]. This edition therefore does *not* identify “one more provability box” with “one more level of self-awareness.”

13.4 A checked self-description scale

Fix a metatheory T capable of coding A , the current search state Σ_m , and the receipts produced by a declared family of self-description checkers. A self-description counts at any level only when its checker verifies the claimed relationship between the descriptor and the actual coded machine/state. The following scale is an architectural convention for this book, not a universal hierarchy of minds.

Definition 13.1 (Formal self-description levels). Relative to a declared architecture and self-description checker family:

- **Level 0:** no nontrivial checked self-description of the designated kind is available to the controller.
- **Level 1:** checked *static* self-description is available: for example, the machine can identify its own code or version, the roles it implements, and the verifier environment it is using.
- **Level 2:** Level 1 is available and checked *dynamic* self-state facts can also be represented and used by the controller, such as the current BANK identifier, selected FUTUREBANK dependencies, verifier-history summaries, resource counters, or whether complete fallback has received its reserved steps.
- **Level n** for finite $n \geq 3$: an implementation-specific extension adds a further bounded layer of checked metacontrol description, such as representing and evaluating how a lower-level controller or self-model would change under candidate search actions. The exact schema must be declared; the symbol n alone does not determine one.
- **Level ω :** every declared finite level schema can in principle be instantiated and checked. This means “all finite levels are supported,” not that the machine stores a completed infinite object or performs infinitely many operations in a finite run.

The scale measures formalized self-description and metacontrol capacity. It does not measure proof depth, intelligence, subjective consciousness, FUTUREBANK size, or chess strength. Level 2 is singled out later only as a plausible *engineering* sweet spot: it exposes useful current-state information to control without automatically incurring the cost of deeper metareasoning. Whether that tradeoff is favorable is an empirical conjecture, not a theorem.

Chapter 14

Imagination and self-awareness are different axes

Mini-roadmap. A chess engine can search deep counterfactual trees without proving any sentence about itself. Conversely, a trivial formal machine can prove a coded statement about itself without possessing a useful FUTUREBANK. We make this separation explicit.

Proposition 14.1 (Separation of basic imagination and basic self-description). *Consider any class of computational search systems containing both (i) a system with nontrivial counterfactual/FUTUREBANK search but no checked Level-1 self-description available to its controller, and (ii) a tagged system with a checked Level-1 self-description but only a constant or nonbranching search policy. Within such a class, nontrivial FUTUREBANK search does not imply Level-1 formal self-description, and Level-1 formal self-description does not imply nontrivial FUTUREBANK search.*

Proof. System (i) is a counterexample to the first implication, and system (ii) is a counterexample to the converse. The claim is therefore a separation result about the two architectural properties, not a claim that either example has rich mathematical intelligence. $\square$

The useful question is therefore not “does self-awareness create imagination?” It does not. The useful question is whether verified self-knowledge can improve the controller’s use of imagination.

14.1 FUTUREBANK-reflective AMLDs

Call an AMLD *FUTUREBANK-reflective* relative to an architecture description if its formal language can name selected search components and it can obtain checker-accepted statements about some of them, such as:

- “my current verified bank is B_m ;
- “branch f depends on speculative lemma λ ;
- “my controller has spent N verifier calls in this quotient region;”
- “these recent candidates were rejected;”

- “complete fallback has not received its scheduled share of steps.”

These are examples of Level-2-style dynamic self-description. They are architecture-relative facts, and they should not be confused with unrestricted claims about the machine’s own soundness or truth.

Chapter 15

Reflection can guide search but cannot redefine proof

Mini-roadmap. We enlarge the full state to include a reflective coordinate. Then a one-line theorem shows why the verifier can remain immune to changes in reflection.

Let

$$\Sigma = (\mathbf{B}, \mathbf{F}, \mathbf{H}, \mathbf{Ref}, \mathbf{Ctrl}).$$

The selector may depend on all coordinates:

$$S : \mathfrak{S} \times \mathbb{N} \longrightarrow \mathcal{K} \cup \{\mathbf{NONE}\}.$$

The intrinsic verifier does not.

Theorem 15.1 (Reflection-neutrality of verification). *Let two search states have the same environment $\mathcal{E}$ and admissible verified bank $\mathbf{B}$ but possibly different *FUTUREBANKs*, histories, controllers, and reflective states. For any fixed semantic certificate $s \in \mathcal{K}_{\text{sem}}$, the intrinsic verifier returns the same acceptance result in both states.*

Proof. This is Theorem 3.6 with the reflective coordinate made explicit. □

Corollary 15.2 (Reflective search without reflective relaxation). *A reflective AMLD may use correct self-descriptive information to alter branch expansion, role allocation, creativity, backtracking, or fallback policy while preserving invariant proof semantics.*

The slogan is:

Formal self-description may tell the machine something about how it is searching. It does not tell the verifier to lower the standard for proof.

Exercise 15.3 (Two axes). Give four hypothetical systems occupying the four cells of a 2×2 table: low/high *FUTUREBANK* capability crossed with low/high formal self-description. Explain why this table is more informative than a single “intelligence level.”

Chapter 16

Safety, liveness, and fair fallback

Mini-roadmap. Verification gives a safety story: false or ill-scoped records should not enter the BANK. It does not by itself give a liveness story: the machine might search the wrong place forever. Fair scheduling and complete fallback address a different question.

16.1 Safety versus liveness

Safety asks: can an invalid candidate enter the verified bank or cause a false terminal status?

Liveness asks: if a discoverable certificate exists, will some component of the system eventually reach it?

A fixed sound verifier strongly supports safety. It does not ensure liveness because a selector can ignore the right proof forever.

16.2 Fair schedules

A schedule is fair to a component if that component receives infinitely many opportunities in an unbounded run. A simple round-robin schedule over P, R, I, C is fair. An adaptive scheduler can also be fair if it reserves nonzero recurring attention for each required process.

16.3 Complete fallback

Suppose there exists a baseline search F that is complete for a certificate class of interest. We may want a learned imaginative controller to dominate most of the computation while occasionally yielding to F .

Theorem 16.1 (Creative conservativity under fair complete fallback). *Suppose a fallback procedure F is complete for finding any terminal certificate in a declared recursively searchable certificate class. Assume that F maintains persistent progress state and that each effective fallback step in the combined AMLD advances the same complete search that F would perform in isolation, or a monotone extension of it that never removes a candidate that the isolated search would eventually examine. If*

the AMLD schedule gives F infinitely many such effective steps, then arbitrary additional heuristic, imaginative, or reflective search interleaved with those steps does not destroy the completeness guarantee supplied by F for that certificate class.

Proof. If a certificate in the covered class exists, completeness of the isolated fallback search implies that it is encountered after some finite number N of effective fallback advances. By the persistence and coverage-preservation hypotheses, interleaving does not reset the fallback state or delete any candidate needed for those N advances. Fair interleaving eventually supplies at least N effective fallback steps. Other search may delay the calendar time of the N th fallback step, but it cannot prevent the covered certificate from being reached. $\square$

This theorem is conditional and modest. It does not make undecidable theories decidable. It says that imaginative search can be added conservatively around a complete procedure for whatever certificate class that procedure already covers.

Part V

Applications and Limits

Chapter 17

Unlimited hypernatural budgets: what they do and do not buy

Roadmap. Corollary 1.9 says that a finite-resource AMLD may honestly return UNKNOWN even when the metatheoretic status of a standard sentence is fixed. This chapter studies an idealized unlimited nonstandard-natural bound. The formal results require only an elementary extension of the standard structure of arithmetic; first-order compactness and the elementary diagram supply such an extension together with an element larger than every standard natural. A richer nonstandard-analysis enlargement is needed only if one wants to speak literally about internal hyperfinite sets, hyperfinite BANKs, or an internal run of hyperfinite length. Keeping those two frameworks separate is essential. We then prove coverage of every standard finite PA proof, derive an external classification of standard PA sentences, and explain why none of this yields an ordinary computable decision procedure or an automatic finite independence certificate.

17.1 The nonstandard setting

Let $\mathbb{N}$ denote the standard first-order structure of arithmetic in the usual language of PA. The arithmetic results below use an elementary extension

$$\mathbb{N} \preceq M.$$

Here $\mathbb{N} \preceq M$ means that the inclusion is elementary: every first-order formula with parameters from the embedded standard copy has the same truth value in $\mathbb{N}$ and in M . We reserve the star notation $^*\mathbb{N}$ for the later, optional full nonstandard-analysis framework. This notational separation prevents an elementary extension of arithmetic from being mistaken for an entire nonstandard universe. We first justify the existence of M with an element above every standard natural.

17.1.1 Compactness supplies the unlimited horizon

The first-order *Compactness Theorem* says that a set T of first-order sentences has a model whenever every finite subset of T has a model. Enderton proves this standard formulation in Section 2.5 after the completeness theorem [14]. Marker treats compactness in Section 2.1 and elementary embeddings and elementary diagrams in Section 2.3 [15].

Expand the language of arithmetic by adding a constant symbol c_a for every $a \in \mathbb{N}$. Let

$$\text{ElDiag}(\mathbb{N})$$

denote the *elementary diagram* of $\mathbb{N}$: all sentences in this expanded language that are true in $\mathbb{N}$ when each c_a is interpreted as a . Add one further constant symbol H , and consider

$$T_H = \text{ElDiag}(\mathbb{N}) \cup \{H > c_n : n \in \mathbb{N}\}.$$

Theorem 17.1 (Compactness existence of an unlimited nonstandard natural). *The theory T_H has a model. Consequently there is a proper elementary extension of the standard natural-number structure containing an element greater than every standard natural number.*

Proof. Take any finite subset $T_0 \subseteq T_H$. Only finitely many inequalities involving H occur in T_0 , say

$$H > c_{n_1}, \dots, H > c_{n_k}.$$

If no such inequality occurs, choose any ordinary interpretation for H . Otherwise interpret every c_a by the corresponding standard natural a and choose an ordinary natural $m > \max\{n_1, \dots, n_k\}$ to interpret H . Every elementary-diagram sentence in T_0 is true by definition, and the finitely many displayed inequalities are true by the choice of m . Thus every finite subset of T_H has a model.

By Compactness, T_H has a model M_H . Since $M_H \models \text{ElDiag}(\mathbb{N})$, the map

$$a \mapsto c_a^{M_H}$$

is an elementary embedding of $\mathbb{N}$ into the arithmetic reduct of M_H . Identify $\mathbb{N}$ with its image and write M for the resulting elementary extension. The interpretation of H satisfies

$$H > n \quad \text{for every standard } n \in \mathbb{N}.$$

Therefore H is outside the embedded standard copy, the extension is proper, and H is an unlimited nonstandard natural. $\square$

Remark 17.2 (Existence is not execution). The theorem is an existence result in first-order model theory. It does not provide a standard natural number that can stand in for H , a canonical computable presentation of M , or an ordinary finite procedure that reaches stage H . Compactness constructs no “fantastically large ordinary integer.”

Remark 17.3 (Why the elementary diagram is used). If one added $H > \bar{n}$ for every standard numeral merely to the complete theory $\text{Th}(\mathbb{N})$, Compactness would give a nonstandard model elementarily equivalent to $\mathbb{N}$. The present chapter wants the more specific relation $\mathbb{N} \preceq M$. Naming every standard element and including the full elementary diagram ensures that the canonical map from the named standard copy into the compactness model is elementary. Marker gives the elementary-diagram machinery in Section 2.3 [15].

Important distinction. The use of $\text{ElDiag}(\mathbb{N})$ is a strong semantic hypothesis, not harmless notation. It packages the full first-order truth of the standard structure with named standard parameters. In particular, the resulting arithmetic extension satisfies $M \models \text{Th}(\mathbb{N})$, not merely $M \models \text{PA}$. Compactness then proves that a suitable elementary extension exists, but it does not give an effective procedure for listing all of $\text{ElDiag}(\mathbb{N})$, computing a presentation of M , or deciding its satisfaction relation. The exact classification results later in the chapter rely on this elementarity; they are not consequences of “being a nonstandard model of PA” alone.

Remark 17.4 (Terminology: nonstandard natural versus hypernatural). The compactness theorem above produces an elementary extension M of arithmetic. Its new elements are unambiguously *nonstandard natural numbers*. In a full nonstandard-analysis framework equipped with a star map or enlargement, elements of ${}^*\mathbb{N}$ are conventionally called *hypernaturals*. This chapter retains “hypernatural” in its title because that is the intended AMLD idealization, but the proofs of the arithmetic results need only M . If a full NSA universe is later chosen, its natural-number sort may play the role of M . Goldblatt gives a standard treatment of transfer, internal and external sets, hyperfinite sets, and nonstandard frameworks, while Keisler provides a complementary infinitesimal-calculus treatment [17, 16].

We now fix one such elementary extension

$$\mathbb{N} \preccurlyeq M$$

and one $H \in M$ satisfying

$$H > n \quad \text{for every standard } n \in \mathbb{N}.$$

We call such an H *unlimited*. In the arithmetic framework, “unlimited” is a metatheoretic condition referring to the embedded standard copy; it is not a single first-order predicate of the language of arithmetic. In a full nonstandard-analysis enlargement, the collection of unlimited hypernaturals is external rather than internal. See Goldblatt’s chapters on internal/external and hyperfinite entities [17].

Checkpoint. The logical dependency is: **Compactness gives existence; the elementary diagram gives the elementary embedding; elementarity transfers first-order statements with standard parameters; and unlimitedness gives coverage of every standard finite proof code.** Elementarity supplies standard correctness; unlimitedness supplies the bound. Neither may be silently substituted for the other, and none of these steps gives a standard computer an executable value of H .

17.1.2 Two frameworks that must not be conflated

There are two legitimate readings of the discussion, but they support different vocabulary.

1. **Arithmetic elementary-extension framework.** Work only with $\mathbb{N} \preccurlyeq M$, the PA proof predicate interpreted in M , and a nonstandard bound $H \in M$. This is sufficient for every theorem in this chapter. One should speak about truth in the elementary extension and bounded quantified formulas, not assume a whole universe of internal sets.
2. **Full nonstandard-analysis framework.** Work in a suitable nonstandard enlargement of a set-theoretic or superstructure universe. Then the natural-number sort contains hypernaturals, the interval $[0, H]$ is an internal hyperfinite set, and one may literally discuss internal functions and hyperfinite runs. Goldblatt treats internal functions and hyperfinite sets explicitly [17].

The second framework is useful for the AMLD picture, but it is stronger infrastructure than the arithmetic compactness theorem itself supplies. No theorem below depends on silently moving from the first framework to the second.

17.1.3 The bounded proof-existence condition

Fix a standard Gödel coding of PA formulas and proofs. For a standard PA sentence φ , let

$$\text{Prf}_{\text{PA}}(p, \ulcorner \varphi \urcorner)$$

be a standard arithmetical formula expressing that p codes a valid PA proof with conclusion code $\ulcorner \varphi \urcorner$. The same formula is interpreted in the elementary extension M .

For an unlimited $H \in M$, define the bounded proof-existence condition

$$B_H^M(\varphi) \quad :\Longleftrightarrow \quad M \models (\exists p)[p \leq H \wedge \text{Prf}_{\text{PA}}(p, \ulcorner \varphi \urcorner)].$$

The superscript records the ambient elementary extension. This definition is model-theoretic: it states that a bounded first-order formula is true in M . It does not presuppose that an external observer can execute a search of length H .

Proposition 17.5 (Why elementarity is essential). *The reverse direction in Theorem 17.6 below depends essentially on $\mathbb{N} \preccurlyeq M$. Replacing M by an arbitrary nonstandard model of PA can make a nonstandard proof witness appear for a standard sentence that has no ordinary PA proof.*

Proof. A standard example comes from the second incompleteness theorem. Assume PA is consistent and use the usual arithmetized consistency sentence $\text{Con}(\text{PA})$. Then PA does not prove $\text{Con}(\text{PA})$; hence PA together with $\neg \text{Con}(\text{PA})$ is consistent and has a model A by completeness. In A , the sentence $\neg \text{Con}(\text{PA})$ asserts that some p satisfies the PA proof predicate for a contradiction. Since PA is externally consistent, no standard natural is such a proof code. In a model of PA the standard numerals form an initial segment of the arithmetic order, so any such witness $p \in A$ is nonstandard and lies above every standard numeral. Taking $H = p$, the analogue of $B_H^A(\perp)$ holds even though there is no ordinary finite PA proof of $\perp$. Enderton treats arithmetized provability and the second incompleteness theorem in Section 3.7 [14]. $\square$

Thus two ingredients have different jobs: the unlimited bound ensures that no *standard* proof code is missed, while elementarity of the extension prevents spurious nonstandard witnesses for false standard proof-existence statements. This is why the chapter constructs an elementary extension of the standard structure rather than merely choosing an arbitrary nonstandard model of PA.

If M is realized as the natural-number sort ${}^*\mathbb{N}$ of a full nonstandard-analysis enlargement, then $B_H^M(\varphi)$ may be pictured as the result of an internal hyperfinite exhaustive scan of

$$[0, H] = \{p \in {}^*\mathbb{N} : 0 \leq p \leq H\}.$$

We may call that optional idealized scan E_H . The scan picture is pedagogically useful, but the formal results below are stated in terms of B_H^M , so they do not depend on importing internal-set language into the arithmetic-only framework.

17.2 Unlimited-bound proof equivalence

Theorem 17.6 (Unlimited-bound proof equivalence). *Let $H \in M$ be unlimited and let φ be a standard PA sentence. Then*

$$\text{PA} \vdash \varphi \quad \Longleftrightarrow \quad B_H^M(\varphi).$$

The same equivalence holds with $\neg\varphi$ in place of φ , and with $\perp$ in place of φ .

Proof. Suppose first that $\text{PA} \vdash \varphi$. Some ordinary finite PA proof of φ has a standard Gödel code $n \in \mathbb{N}$. Since H is unlimited, $n < H$. The elementary embedding preserves the standard proof predicate on the standard parameter $\ulcorner \varphi \urcorner$, so the same standard code witnesses $B_H^M(\varphi)$.

Conversely, suppose $B_H^M(\varphi)$. Then in particular

$$M \models (\exists p) \text{Prf}_{\text{PA}}(p, \ulcorner \varphi \urcorner).$$

The formula defining the PA proof predicate and the code $\ulcorner \varphi \urcorner$ are standard. Since $\mathbb{N} \preccurlyeq M$, elementarity gives

$$\mathbb{N} \models (\exists n) \text{Prf}_{\text{PA}}(n, \ulcorner \varphi \urcorner).$$

Hence an ordinary finite PA proof of φ exists. The bound $p \leq H$ is unnecessary in this reverse direction once any proof witness exists in the elementary extension. $\square$

Corollary 17.7 (Unlimited-bound coverage of standard positive certificates). *Every standard finite P-, R-, or C-certificate for PA has a standard code below every unlimited H . Thus the bounded condition B_H^M cannot miss an existing standard PA proof of φ , $\neg\varphi$, or $\perp$.*

Corollary 17.8 (Standard positive-witness existence). *If $B_H^M(\varphi)$ holds for a standard PA sentence φ , then an ordinary finite PA proof of φ exists. A witness to the bounded formula in M need not itself be standard.*

The last sentence matters. Elementarity proves existence of some standard proof; it is not an algorithm that extracts that proof from a particular nonstandard witness. Once standard existence is known metatheoretically, ordinary enumeration of finite PA proofs eventually finds a standard witness, but no effective shortcut has been obtained merely from the nonstandard one.

17.3 External classification, not an ordinary decider

Assume for this section that PA is consistent. For each standard PA sentence φ , define the metatheoretic classification

$$D_H^M(\varphi) = \begin{cases} P, & \text{if } B_H^M(\varphi), \\ R, & \text{if } B_H^M(\neg\varphi), \\ I, & \text{if neither bounded condition holds.} \end{cases}$$

Here I means exactly

$$\text{PA} \not\vdash \varphi \quad \text{and} \quad \text{PA} \not\vdash \neg\varphi.$$

This is the derivability-based meaning of “independent of PA” used in this chapter. If consistency is not assumed, a C-condition for $B_H^M(\perp)$ should receive the contradiction-dominance priority described in Theorem 1.5.

Theorem 17.9 (Unlimited-bound trichotomy for standard PA sentences). *Let PA be consistent, let $H \in M$ be unlimited, and let φ be a standard PA sentence. Exactly one of the following holds:*

- (i) $B_H^M(\varphi)$;
- (ii) $B_H^M(\neg\varphi)$;

(iii) neither $B_H^M(\varphi)$ nor $B_H^M(\neg\varphi)$.

Moreover these cases correspond respectively to

$$\text{PA} \vdash \varphi, \quad \text{PA} \vdash \neg\varphi, \quad \text{PA} \not\vdash \varphi \text{ and } \text{PA} \not\vdash \neg\varphi.$$

Proof. The equivalences are immediate from Theorem 17.6. Consistency excludes simultaneous occurrence of the first two cases, because proofs of both φ and $\neg\varphi$ yield inconsistency in ordinary classical PA. The third case is exactly the failure of both derivability statements. $\square$

Proposition 17.10 (Model-and-cutoff invariance on standard targets). *Let $\mathbb{N} \preceq M_1$ and $\mathbb{N} \preceq M_2$ be elementary extensions, and let $H_1 \in M_1$ and $H_2 \in M_2$ each be larger than every standard natural. For every standard PA sentence φ ,*

$$D_{H_1}^{M_1}(\varphi) = D_{H_2}^{M_2}(\varphi).$$

Thus, on standard targets, the subscript and superscript record the idealization but do not change the classification.

Proof. By Theorem 17.6, each bounded positive condition is equivalent to the same standard derivability statement $\text{PA} \vdash \varphi$ or $\text{PA} \vdash \neg\varphi$. Therefore the classification cannot depend on which elementary extension or which unlimited cutoff was chosen. $\square$

Important distinction. The map D_H^M is a metatheoretically defined classification of standard syntax, not an effective procedure presented to a standard computer. In the full nonstandard-analysis reading, an internal hyperfinite scan E_H can represent the bounded search, but a standard observer is not thereby given a way to run to its last stage or read its result as ordinary algorithmic output. In the arithmetic-only reading, it is cleaner to regard B_H^M simply as a truth condition in the chosen elementary extension.

Corollary 17.11 (Idealized external settlement). *Under the hypotheses of Theorem 17.9, D_H^M assigns every standard PA sentence its correct P , R , or I metatheoretic case. If a C -condition is included and PA is inconsistent, every standard proof of contradiction is likewise below every unlimited cutoff.*

This result does not contradict undecidability of ordinary computable PA theoremhood. Neither Compactness nor the bounded truth condition B_H^M supplies a standard program with a standard finite input and standard terminating runtime. The idealization uses a nonstandard bound together with the external/metatheoretic fact that the bound lies above every standard natural.

17.4 Why no ordinary finite cutoff can replace an unlimited one

Theorem 17.12 (Finite-cutoff insufficiency). *For every standard $N \in \mathbb{N}$, there exists a standard PA theorem θ for which no proof code of θ is at most N .*

Proof. There are only finitely many natural numbers $p \leq N$, hence only finitely many proof codes at most N , and therefore only finitely many theorem conclusions represented by those codes. PA has infinitely many distinct theorems; for example, the reflexive equalities for distinct numerals give infinitely many standard theorem statements. Consequently at least one PA theorem has no proof whose code is at most N . $\square$

Corollary 17.13 (No universal standard exhaustion bound). *Failure to find a proof below one fixed standard finite cutoff N cannot, by itself, establish nonderivability for arbitrary PA targets. No standard N has the cofinal coverage property of an unlimited H .*

This theorem makes precise why “replace H by a gigantic ordinary integer” does not preserve the argument. However large the ordinary integer is, some ordinary finite PA proofs lie beyond it.

17.5 The direct-standardization barrier

Proposition 17.14 (Direct-cutoff standardization barrier). *Let $H \in M$ be unlimited. There is no standard natural N that can replace H while preserving the property*

$$\text{every standard } n \in \mathbb{N} \text{ satisfies } n \leq H.$$

Consequently a bounded nonfinding statement such as

$$M \models (\forall p \leq H) \neg \text{Prf}_{\text{PA}}(p, \ulcorner \varphi \urcorner)$$

cannot be turned into an equally exhaustive ordinary finite-search record merely by substituting a standard cutoff.

Proof. If a standard N dominated every standard natural, then the standard natural $N + 1$ would satisfy $N + 1 > N$, a contradiction. Replacing H by any standard N therefore discards standard proof codes larger than N . $\square$

Important distinction. No “standard-part” operation repairs this problem. In a full hyperreal framework, standard part is defined for finite hyperreals, not as a natural-valued operation sending an unlimited hypernatural to an ordinary cutoff with the same cofinal property; see Goldblatt [17]. In the P , R , and C cases, Theorem 17.6 gives existence of an ordinary finite proof witness. In the I case, bounded nonfinding together with unlimitedness gives the external nonderivability conclusion, but it does not automatically produce a finite certificate in the AMLD’s declared $\mathcal{K}_I$ format.

No conservativity or proof-translation theorem for a particular formalization of nonstandard analysis is assumed here. If some chosen formal NSA system is separately known to be conservative over, or interpretable in, an appropriate standard metatheory, that is an additional theorem about those formal systems. It should not be confused with replacing the particular nonstandard cutoff H by a standard finite cutoff.

17.6 Three conclusions to keep separate

One preliminary fact comes first: Compactness plus the elementary diagram supplies a proper elementary extension containing an unlimited nonstandard natural. Once such an environment is fixed, the search-bound argument establishes three further points:

1. **Coverage.** Every standard finite PA proof code lies below every unlimited cutoff.

2. **External classification.** Under elementarity and the stated consistency assumption, the bounded conditions give the correct external $P/R/I$ classification of each standard PA sentence.
3. **Ordinary computability.** No ordinary algorithm deciding PA has been obtained. An unlimited H is not a standard finite resource bound, and the negative I -case is not automatically a standard finite AMLD certificate.

The optional full-NSA picture adds a literal internal hyperfinite scan, but it does not alter any of these conclusions.

Corollary 17.15 (Reconciliation with computational UNKNOWN). *Corollary 1.9 and Theorem 17.9 are compatible. An ordinary AMLD with a standard finite budget may return UNKNOWN. The nonstandard construction supplies an external classification theorem for standard PA sentences under the hypotheses above. It does not turn every external I -classification into a standard finite I -certificate, and it does not make PA decidable by an ordinary program.*

Exercise 17.16 (Building the unlimited horizon by compactness). Let $\mathcal{L}(\mathbb{N}, H)$ be the language of arithmetic expanded by constants c_a for each $a \in \mathbb{N}$ and one additional constant H . Let

$$T_H = \text{ElDiag}(\mathbb{N}) \cup \{H > c_n : n \in \mathbb{N}\}.$$

- (a) Explain why every finite subset of T_H is satisfiable in the standard natural numbers after choosing a sufficiently large ordinary interpretation for H .
- (b) Apply Compactness to conclude that T_H has a model.
- (c) Explain why the named copy of $\mathbb{N}$ embeds elementarily into that model and why the interpretation of H is outside the standard copy.
- (d) Explain why the definition of “ H is larger than every standard natural” refers to the externally designated standard copy and is presented here as an infinite family of inequalities, not as an extra predicate added to PA.
- (e) Explain why the existence proof gives no ordinary finite algorithm that counts to H .

Exercise 17.17 (Arithmetic core versus full NSA). Which parts of the chapter require only an elementary extension of arithmetic, and which phrases require the richer full nonstandard-analysis framework? In particular, classify each phrase: (a) $B_H^M(\varphi)$; (b) “internal hyperfinite interval”; (c) “elementarity transfers the proof-existence statement”; (d) “hyperfinite BANK history.”

Exercise 17.18 (Why an arbitrary nonstandard PA model is not enough). Assume PA is consistent. Explain why a model of $\text{PA} + \neg \text{Con}(\text{PA})$ can contain a nonstandard number that the model regards as a PA proof code for contradiction. Which step of Theorem 17.6 fails if one works in this model instead of an elementary extension of the standard $\mathbb{N}$?

Exercise 17.19 (Unlimited versus merely enormous). Let $H \in M$ be unlimited and let $N \in \mathbb{N}$ be an extremely large standard integer.

- (a) Why does mere membership $H \in M$ not by itself state unlimitedness?
- (b) If PA proves a standard sentence φ , why must $B_H^M(\varphi)$ hold?
- (c) Why can no fixed standard N replace H in that universal coverage statement?
- (d) Explain the asymmetry between a positive $P/R/C$ witness and the negative I -classification.
- (e) Why does the existence of the external classification D_H^M not contradict undecidability of ordinary computable PA theoremhood?

Chapter 18

Authentication as one worked special case

Roadmap. This chapter is an application, not the destination of the theory. Passwords and authentication are used because they sharply separate *search*, *verification*, *available information*, and *system design*. We will model a password verifier, define “flood control” as a hard or effective bound on online attempts, prove several finite-attempt theorems, distinguish online from offline verification, and show why a more plausible defensive use of AMLD-style reasoning is often to settle formal claims about an authentication system rather than to “guess the password.” The same AMLD architecture is intended for much broader application classes, including ordinary mathematical settlement, software and protocol verification, engineering safety, and formalized scientific reasoning. The discussion here is about authorized analysis, defensive design, and formal modeling.

18.1 Why passwords belong in this book

A password system looks superficially like proof search. A candidate string is proposed; a verifier returns an answer; a search policy decides what to try next. That resemblance is real, but it must not be exaggerated.

A mathematical verifier answers a question such as “is this certificate a valid proof in environment $\mathcal{E}$?” A live authentication verifier is usually *stateful*: it may count failures, insert delays, require an additional factor, reject an otherwise correct credential after a lockout, or stop accepting attempts altogether. For that reason, password authentication is best treated as a *special case of guarded search with verification*, not as literal mathematical theorem proving.

This chapter also illustrates a broader policy point. Powerful search machinery should be developed together with narrow correctness boundaries and explicit use constraints. In the companion essay *Fund the Proof Machines—Guard the Capabilities That Can Become Weapons*, Tenneson argues that automated theorem proving differs from human-mimicking AI because its distinctive power is the search for independently checkable chains of reasoning. The essay treats passwords as only one dual-use example alongside mathematical discovery, software verification, protocol analysis, engineering safety, cryptographic analysis, and other formally stated targets [10]. That is also the role passwords play here: they are a revealing example, not the organizing purpose of AMLD.

Important distinction. Nothing in the AMLD architecture acts like a magic eight ball that can pluck a secret password from the sky. If a password is independent of everything the machine can observe, then imagination, reflection, a larger FUTUREBANK, or a more elaborate COMPASS cannot manufacture information about that secret. They can only change the order in which candidates are considered. Theorems below make this precise.

18.2 A minimal authentication model

Let S be a finite set of candidate secrets and let the actual secret be a random variable $W \in S$. Let Z denote all side information legitimately available to the analysis. For the realized side-information state z , assume the model supplies the conditional distribution

$$\pi_z(w) = \Pr(W = w \mid Z = z).$$

When z is fixed, abbreviate π_z by π . This is the search prior for the modeled episode. A secret generated uniformly and independently of Z makes π uniform on the generator's support. Human choice, reuse, policy constraints, previously disclosed information, or implementation artifacts may instead make π highly nonuniform.

For a live system, let

$$\text{Auth}_t(g, x_t) \in \{\text{ACCEPT}, \text{REJECT}, \text{DELAY}, \text{BLOCKED}\}$$

be the authentication response to candidate g in system state x_t . The state may include a retry counter, lockout timer, risk score, second-factor state, account state, and other policy variables.

Definition 18.1 (Flood control). In this book, *flood control* is an umbrella term for controls that bound or slow repeated authentication attempts. Standard security terminology includes *rate limiting*, *login throttling*, and *account lockout*. If, within the explicitly modeled episode or time horizon, the system will process at most N candidate submissions, we call N the *online attempt budget* for that episode.

A delay schedule by itself does not automatically create a finite N : if the modeled horizon is unbounded, delayed attempts may continue indefinitely. The finite-attempt theorems below apply only when the episode definition, lockout rule, disabling rule, or finite time horizon supplies an actual or effective cap on the number of processed guesses.

For centrally verified passwords, NIST SP 800-63B-4 requires a rate-limiting mechanism. Unless an authenticator-specific rule says otherwise, its general rule limits consecutive failed attempts using a specific authenticator on one subscriber account to no more than 100 by disabling that authenticator; 100 is an upper bound, and lower limits are permitted [11]. OWASP likewise recommends login throttling and discusses lockout thresholds, observation windows, lockout duration, and the denial-of-service risk of overly rigid lockout policies [12].

18.3 The right quantity is top- N probability mass

Suppose the candidates are ordered so that

$$\pi_{(1)} \geq \pi_{(2)} \geq \cdots \geq \pi_{(|S|)}.$$

Define the *top- N exposure*

$$G_N(\pi) = \sum_{i=1}^{\min(N, |S|)} \pi_{(i)}.$$

This quantity answers a concrete question: if only N online guesses are possible, what is the greatest success probability available to a strategy that receives no extra information except ordinary accept/reject responses?

Theorem 18.2 (Optimal finite-attempt guessing bound). *Assume a strategy may submit at most N distinct candidates, receives no side channel or new information about W other than whether a submitted candidate equals the secret, and cannot bypass the attempt budget. Then its maximum probability of success is exactly*

$$G_N(\pi).$$

An optimal strategy tries candidates in nonincreasing order of π .

Proof. Any strategy that makes at most N distinct guesses succeeds only when W lies in the set of guessed candidates. The probability mass of any such set is at most the sum of the N largest point probabilities, namely $G_N(\pi)$. Guessing those N candidates attains that bound. Rejections merely remove already-tried candidates; under the stated assumptions they reveal no additional information that changes the relative ordering of the untried prior masses. $\square$

Corollary 18.3 (Uniform-secret flood-control bound). *If W is uniform on S , then any such N -attempt online strategy succeeds with probability at most*

$$\min\left(1, \frac{N}{|S|}\right).$$

This is a clean mathematical form of the motivating intuition that flood control can make an otherwise powerful search method nearly irrelevant to direct online password guessing. If $N \ll M$ and the prior is close to uniform, there simply is not enough interaction budget for sophisticated search to matter much.

18.4 No-information imagination cannot create information

The AMLD vocabulary lets us say the same thing more structurally.

Theorem 18.4 (No-oracle theorem for password imagination). *Under the hypotheses of Theorem 18.2, adding a FUTUREBANK, a COMPASS, self-description, multiple search agents, or arbitrary internal computation cannot raise the success probability above $G_N(\pi)$ unless the added mechanism has access to information that changes the conditional distribution of W or changes the verifier/attempt constraints.*

Proof. Internal computation can choose an ordering or a set of at most N candidates, but success still implies that W lies in that set. Theorem 18.2 bounds the probability mass of every such set. To exceed the bound, some hypothesis must change: new information must alter π , the verifier must leak information, the candidate equivalence relation must differ from the model, or the attempt policy must be bypassed or altered. $\square$

Design principle. For passwords, imagination is therefore best understood as *hypothesis generation about structure*: candidate distributions, verifier behavior, recovery routes, storage assumptions, policy inconsistencies, or specification defects. It is not a source of secret bits.

18.5 A min-entropy version of the flood-control theorem

The earlier distribution π is already conditioned on the side information available for the episode. To avoid ambiguity with several inequivalent notions of conditional min-entropy used in information theory, define the min-entropy of this *realized conditional distribution* by

$$h_\infty(\pi) = -\log_2 \left(\max_w \pi(w) \right).$$

Theorem 18.5 (Flood-control min-entropy bound). *Under the hypotheses of Theorem 18.2,*

$$\Pr(\text{success in at most } N \text{ guesses} \mid Z = z) \leq \min \left(1, N 2^{-h_\infty(\pi)} \right),$$

where z is the fixed realized side information used to define π .

Proof. Every point probability is at most $2^{-h_\infty(\pi)}$. The probability mass of any N guessed candidates is therefore at most N times that quantity, capped by 1. $\square$

This theorem explains why a small N is powerful only when it is combined with a password distribution whose largest point probabilities are themselves small. The nominal size of the character alphabet is not enough. What matters to a capped online attacker is how much probability mass lies in the first few guesses.

18.6 Human-chosen passwords, random passwords, and password managers

The top- N theorem turns password quality into a distributional statement.

- A unique password drawn from a high-entropy generator can make π close to flat over a large effective support, keeping $G_N(\pi)$ small for realistic N .
- Reuse or highly predictable human choice concentrates π , making a small set of candidates carry more probability mass.
- Password blocklists remove especially high-probability choices from the admissible set. NIST explicitly connects blocklists to the fact that online attempts are rate-limited [11].
- Password managers can generate and store long, unique secrets, reducing the human incentive to reuse memorable strings.

Thus “strong password” should not be reduced to a composition-rule slogan. In the finite-attempt model, the central operational quantity is $G_N(\pi)$ or a bound on it.

18.7 Online and offline verification are different worlds

Flood control belongs to the *live verifier*. It may be decisive online and irrelevant offline.

Proposition 18.6 (Online/offline separation). *Suppose an online service enforces an attempt budget N , but a compromise exposes data sufficient to test password guesses locally without querying that service. Then the online attempt cap does not bound the number of local verification trials. The defensive cost per offline trial and the password distribution become separate controlling factors.*

Proof. The online cap is enforced by the live service state. A local verifier that does not consult that state does not consume the service’s retry counter. Therefore the online bound does not apply to those local evaluations. $\square$

This is why password storage matters. OWASP recommends one-way password hashing rather than reversible encryption, unique salts, and deliberately expensive password-hashing functions so that each offline candidate test has nontrivial cost [13].

Proposition 18.7 (Per-salt precomputation separation). *Model a stored verifier as $y = F(s, w)$ where s is an account-specific salt. Suppose a precomputed table contains only pairs of the form $(w, F(s_1, w))$. For a different salt s_2 , that table does not by itself provide the values $F(s_2, w)$ unless one also has a method for deriving or computing the s_2 -parameterized outputs. Thus a verifier for the second account must still use the second salt.*

Proof. The table records outputs only for inputs whose first coordinate is s_1 . An offline check for the second account requires outputs whose first coordinate is s_2 . Without additional information or a derivation that supplies those outputs, the first table alone does not contain them. $\square$

The proposition is intentionally modest. It does not claim that salting makes weak passwords safe or prevents per-account guessing. It states the narrower reason unique salts defeat direct cross-account reuse of a single one-salt password-to-hash lookup table.

18.8 Multi-factor authentication changes the target predicate

A modern authentication decision may not be “password correct?” It may be

$$\text{PasswordOK} \wedge \text{SecondFactorOK} \wedge \text{AccountEligible} \wedge \text{PolicyAllows}.$$

Theorem 18.8 (Conjunctive-factor containment). *If a system accepts only when both a password predicate Pwd and an additional-factor predicate Factor hold, then the set of fully accepted authentication events is a subset of the events satisfying Pwd alone. Consequently, success at the password subproblem is not sufficient for full authentication.*

Proof. Every event satisfying $\text{Pwd} \wedge \text{Factor}$ satisfies Pwd . The converse need not hold. $\square$

OWASP describes multi-factor authentication as one of the strongest defenses against password-related attacks [12]. In AMLD terms, this means that the terminal certificate predicate has changed: a password candidate is no longer the whole certificate.

18.9 The more interesting target: prove the verifier conforms to its specification

Direct secret guessing is often the least interesting formal problem. A much richer defensive target is a claim about the authentication system itself.

Let $\text{Req}(z)$ be the *required/intended* acceptance predicate for a complete authentication event z , and let $\text{Impl}(z)$ be the implemented acceptance predicate. For the security question of unintended acceptance, consider the one-sided conformance (acceptance-soundness) claim

$$c_{\text{auth}} : \quad \forall z [\text{Impl}(z) \Rightarrow \text{Req}(z)].$$

This says: every event the implementation accepts is an event the security specification intends to accept. It deliberately does *not* say that every intended event is actually accepted. Full extensional conformance would additionally require the reverse implication, equivalently

$$\forall z [\text{Impl}(z) \leftrightarrow \text{Req}(z)],$$

and failures of that reverse direction represent false rejection, availability, or completeness defects rather than unintended acceptance. The one-sided target is used below because it isolates the safety property relevant to an unauthorized-acceptance witness. The mnemonic predicate names **Req** and **Impl** also avoid reusing the book's bare one-letter role symbols.

Theorem 18.9 (Authentication-defect certificate). *Assume that, for a concrete event z^* , the declared audit checker can verify both $\text{Impl}(z^*) = \text{true}$ and $\text{Req}(z^*) = \text{false}$ from finite representations. Then z^* together with those checker receipts is a finite refutation certificate for*

$$c_{\text{auth}} : \quad \forall z [\text{Impl}(z) \Rightarrow \text{Req}(z)].$$

Proof. The verified witness z^* is a counterexample to the universal implication $\forall z [\text{Impl}(z) \Rightarrow \text{Req}(z)]$. The finite checker receipts make the counterexample independently checkable in the declared audit environment. $\square$

This is the sense in which AMLD-like reasoning may be more useful for passwords: not as a machine that guesses secrets, but as a machine that searches for *certifiable flaws in the modeled authentication system*. Examples of defensively relevant specification surfaces include retry policies, recovery flows, inconsistent enforcement of second factors, account-state transitions, password-storage configuration, overly informative error behavior, parser/normalization consistency, and mismatches between documented policy and implemented policy. OWASP specifically warns that authentication error behavior can disclose account-existence information and recommends generic responses [12].

Important distinction. A defect witness is not a recovered password. The two are different mathematical objects. The useful security question is often “does the implementation admit any event forbidden by the specification?” rather than “what is the secret?”

18.10 A four-role AMLD audit of an authentication system

This one-sided conformance target gives the four AMLD roles concrete defensive meanings.

P : prove c_{auth} for the modeled state space;

R : produce a checked counterexample event z^* ;

I : certify, in the declared metatheory, that c_{auth} is independent of the admitted audit theory;

C : derive an inconsistency in the declared security specification or model assumptions.

The BANK begins with the formal authentication specification, the modeled implementation semantics, and explicitly admitted environmental assumptions. Verified invariants and checked test lemmas can be deposited as the audit proceeds. The FUTUREBANK may hold possible counterexample traces or repair hypotheses, but none becomes an established defect or theorem without a checker-accepted certificate.

This is a much better fit to the book’s architecture than pretending that a hidden password is a theorem waiting to be divined.

18.11 Recovery routes and the weakest-route upper bound

Authentication rarely has only one route. There may be password login, password reset, account recovery, delegated identity, support-assisted recovery, remembered sessions, or hardware-backed authentication.

Theorem 18.10 (Weakest-route upper bound). *Suppose any one of routes $\rho_1, \dots, \rho_k$ is sufficient to obtain the same protected account capability, and let C_i be the minimum cost of successfully traversing route ρ_i under a fixed threat model. Then the minimum cost of obtaining the capability satisfies*

$$C_{\text{system}} \leq \min_i C_i.$$

Proof. An actor may choose the least-cost sufficient route. Therefore system-level cost cannot exceed the cost of the easiest sufficient route. $\square$

The theorem is deliberately an upper bound, not a complete security metric. Its lesson is architectural: strengthening the password path while leaving a weaker recovery path unchanged may not strengthen the whole account by the same amount.

18.12 Flood control can itself create a control problem

A hard lockout reduces online guessing but may also make availability vulnerable if arbitrary outsiders can consume the failure budget of someone else’s account. OWASP explicitly notes this denial-of-service tradeoff and discusses alternatives such as delays and graduated controls [12].

This creates an AMLD-style multiobjective control problem. The controller is not optimizing only

minimize unauthorized acceptance.

It is balancing at least

unauthorized acceptance risk, legitimate-user lockout risk,
latency, recovery burden, operational cost.

There is no universal theorem that one fixed N is optimal for every system. The correct N , delay schedule, and recovery policy depend on the threat model and operational environment.

18.13 The stateful-verifier mismatch with mathematical proof

Proposition 18.11 (Authentication acceptance can be nonmonotone in time). *There exist ordinary authentication policies for which the same credential submission is accepted in one system state and blocked in another, even though the credential string itself has not changed.*

Proof. A lockout or disabled-account state provides an immediate example: the underlying password may be correct while policy blocks authentication in the later state. $\square$

This is unlike the intended verifier semantics of the BANK. Once a proof certificate is valid relative to a fixed environment version, later search history should not make that same certificate invalid. The password analogy therefore teaches an important boundary: search architecture may transfer across domains even when verification semantics do not.

18.14 What an AMLD can realistically contribute to password security

A disciplined research program would emphasize the following defensive tasks:

- formalize the intended authentication policy and prove invariants about it;
- search for finite counterexample traces to policy conformance;
- verify that flood-control rules actually impose the modeled attempt budget;
- check that alternate login and recovery routes satisfy the same intended security requirements;
- reason about the distributional quantity $G_N(\pi)$ rather than relying on nominal password-space size alone;
- compare fixed lockout, delay, and adaptive throttling policies under explicit availability constraints;
- verify that password storage follows a declared modern hashing-and-salting policy;
- test whether error channels reveal facts that the formal policy intends to conceal;
- preserve failed audit hypotheses in verifier history without promoting them into the BANK as defects;
- use the FUTUREBANK for hypothetical repairs, then require regression certificates before accepting a repair as safe.

The likely payoff is therefore *finding and certifying flaws in a system or certifying that a modeled class of flaws is absent*, not mysteriously recovering an otherwise information-theoretically hidden password.

18.15 Worked numerical example: when N makes guessing negligible

Suppose a password is uniformly random over a set S with $|S| = 2^{80}$ possibilities and the live verifier will process at most $N = 100$ candidate submissions in the modeled episode. Corollary 18.3 gives

$$\Pr(\text{success}) \leq \frac{100}{2^{80}} \approx 8.27 \times 10^{-23}.$$

No clever ordering changes this because all candidates have equal prior probability. A COMPASS may be computationally impressive and still be informationally irrelevant.

Now change only the prior. If human behavior or leaked side information concentrates, say, 5% of the conditional probability mass into the first 100 candidates, then $G_{100}(\pi) = 0.05$ and the same flood-control value $N = 100$ no longer gives the same protection. The difference is not “more imagination.” It is a different information model.

18.16 Exercises

Exercise 18.12 (Top- N exposure). Let $S = \{a, b, c, d, e\}$ with probabilities

$$(0.40, 0.25, 0.15, 0.10, 0.10).$$

Compute G_1 , G_2 , and G_3 . Then explain why, with no new side information, a more elaborate search architecture cannot exceed those success probabilities under attempt budgets 1, 2, 3.

Exercise 18.13 (Flood control versus offline verification). Explain why a server-side $N = 10$ lockout can strongly constrain online guessing while placing no direct ten-trial limit on a hypothetical local verifier obtained from compromised password-verifier data. Name the separate defense that becomes important in the local-verification setting.

Exercise 18.14 (Turn security into a settlement problem). Write a conformance statement c for a toy authentication system in which acceptance is intended to require both a password predicate $\text{Pwd}(z)$ and second-factor predicate $\text{Factor}(z)$. What kind of object would let R refute c ?

Exercise 18.15 (Flood-control tradeoff). Give a toy example in which lowering the maximum failure count decreases online guessing opportunity but increases the possibility that a legitimate user is denied service. Explain why this is a control-design problem rather than a theorem that “smaller N is always better.”

Part VI

Implementation

Chapter 19

Data structures

Roadmap. The mathematics now becomes software. This Part gives pseudocode for the AMLD state, the narrow verifier interface, FUTUREBANK generation, role scheduling, backfilling, safe quotienting, and invariant tests.

Listing 19.1: Core records and state

```

record VerificationEnvironment:
  gamma_records      # admitted axioms/hypotheses Gamma
  base_scope         # logical scope of the seed assumptions
  target             # c
  language_id
  inference_system_id
  checker_versions
  environment_id

record BankRecord:
  id                 # stable BANK record identifier
  claim
  semantic_kind      # AXIOM, HYPOTHESIS, INTERMEDIATE_THEOREM,
                     # TARGET_PROOF, NEGATED_TARGET_PROOF,
                     # INDEPENDENCE, CONTRADICTION, other declared kinds
  origin             # ENV, IMPORT, P, R, I, or C
  scope
  certificate
  dependency_ids
  verifier_receipt
  source_proposal_id # NONE for ENV/IMPORT seed material
  source_future_node_id # NONE unless accepted from a FUTUREBANK node
  deposited_at
  environment_id

record Proposal:
  id                 # stable proposal identifier for history/audit
  submitting_role    # exactly one of P, R, I, C
  semantic_kind      # what kind of result is claimed
  claim              # formula/metaclaim claimed by the certificate
  certificate_type    # selects the formal checker
  certificate
  scope
  dependency_ids
  source_future_node_id # NONE unless materialized from a FUTUREBANK node

```

```

    environment_id

record LogicalBank:
    records_by_id
    origin_index      # ENV, IMPORT, P, R, I, C -> record IDs
    semantic_kind_index # semantic kind -> record IDs
    dependency_graph

record FutureNode:
    id                # stable FUTUREBANK history-node identifier
    proof_state
    parent_id         # NONE for root; otherwise one realized history parent
    generating_action
    role
    assumed_dependencies
    bank_dependencies
    depth
    trust_grade       # SPEC or LOCAL
    compass_score
    estimated_cost
    attempted = false

record SearchState:
    bank
    futurebank
    history
    reflection_state
    controller_state
    step

```

The core pseudocode treats a **FutureNode** as a *history node*: the root represents the current initial proof state and has no parent or generating action; every nonroot history node has one parent and one generating action. BANK records, proposals, and FUTUREBANK history nodes have stable identifiers because dependencies, history events, and invariant tests refer to them by ID. Every proposal-construction routine is required to allocate a fresh proposal ID. When SELECT materializes a FUTUREBANK history node, it marks that exact node as attempted and copies its ID into the proposal; if the proposal is accepted, UPDATE carries the same source-node ID into the resulting BANK record. BANK record identity is deliberately finer than formula identity: two separately accepted certificates may yield distinct BANK records even when their claims are syntactically identical. Retrieval indexes may coalesce or canonicalize equivalent claims for efficiency, but the underlying record ledger retains the separate receipts and provenance links. Two histories that reach the same visible proof state therefore remain distinct at this layer, so transition scores and provenance are unambiguous. A later implementation may share downstream computation across future-equivalent histories, but only through the safe quotient discipline of Chapter 23; path multiplicity or ancestry that remains task-relevant must be retained separately.

19.1 Seeding from Γ

Listing 19.2: The BANK starts from Gamma

```

function SEED_BANK(env):
    B = empty LogicalBank
    for premise in env.gamma_records:

```

```
    r = BankRecord(
        id = fresh_record_id(),
        claim = premise.formula,
        semantic_kind = premise.kind, # AXIOM or HYPOTHESIS
        origin = ENV,
        scope = premise.scope or env.base_scope,
        certificate = premise.admission_record,
        dependency_ids = {},
        verifier_receipt = "ADMITTED-BY-ENVIRONMENT",
        source_proposal_id = NONE,
        source_future_node_id = NONE,
        deposited_at = 0,
        environment_id = env.environment_id
    )
    B.add(r)
return B
```

This function makes the conceptual starting point explicit: Γ is not hidden in global code. It becomes auditable initial mathematical state. The seed records use origin `ENV`, not a fictitious *P*, *R*, *I*, or *C* origin. Imported certified libraries can use `IMPORT`.

Chapter 20

The narrow verifier and update rule

Mini-roadmap. The most important software interface is intentionally boring. The verifier receives the environment, admissible bank view, and certificate. It does not receive a compass score or creativity parameter.

Listing 20.1: Narrow verification interface

```
function VERIFY(env, bank, proposal):
    assert proposal.submitting_role in {P, R, I, C}
    assert proposal.environment_id == env.environment_id
    deps = bank.resolve_admissible(proposal.dependency_ids,
                                   scope=proposal.scope)
    checker = env.checker_versions[proposal.certificate_type]
    # submitting_role is provenance; it is not a truth input.
    # claim and scope are untrusted proposal fields until checked.
    return checker.check(env,
                        semantic_kind=proposal.semantic_kind,
                        claim=proposal.claim,
                        scope=proposal.scope,
                        dependencies=deps,
                        certificate=proposal.certificate)
```

Listing 20.2: Verifier-gated update

```
function UPDATE(state, proposal, result):
    state.history.append(event(state.step, proposal, result))

    if result != ACCEPT:
        return state, NONE

    # ACCEPT is returned only after the checker validates the semantic kind,
    # claim, scope, dependencies, certificate, and environment. The deposited
    # record copies fields that have therefore passed the declared checker.
    record = make_verified_record(proposal, result, state.step)
    record.id = fresh_record_id()
    record.source_proposal_id = proposal.id
    record.source_future_node_id = proposal.source_future_node_id
    state.bank.add(record)
    if proposal.source_future_node_id is not NONE:
        state.futurebank.mark_promoted(proposal.source_future_node_id,
                                       bank_record_id=record.id)
```

```
return state, record
```

20.1 Terminal classification

Listing 20.3: Terminal settlement

```
function TERMINAL_STATUS(env, bank, accepted_record):
    # If inconsistency is already certified, it dominates in explosive logic.
    if bank.has_verified_contradiction(env):
        return INCONSISTENT
    # Otherwise classify this accepted terminal record by verified semantics,
    # not by who found it.
    if accepted_record.semantic_kind == TARGET_PROOF \
        and proves_target(env.target, accepted_record):
        return PROVED
    if accepted_record.semantic_kind == NEGATED_TARGET_PROOF \
        and proves_negated_target(env.target, accepted_record):
        return REFUTED
    if accepted_record.semantic_kind == INDEPENDENCE \
        and proves_independence(env.target, accepted_record):
        return INDEPENDENT
    return NONE
```

The pseudocode assumes explosive logic and gives an already verified contradiction certificate priority over other terminal evidence. This is an operational precedence rule, not clairvoyance: if a proof of c is accepted and the run halts before any contradiction is discovered, the machine cannot retroactively know that a future C -search would have succeeded. In a different logic, terminal precedence belongs in the verification environment. The classifier deliberately does not test the record's origin. If R submits an accepted target-proof certificate, the provenance remains R while the semantic status is **PROVED**. Also, first-certificate reporting should not be confused with a proof that the four evidence conditions are globally exclusive; Corollary 1.6 and Theorem 1.7 state the relevant distinction.

Chapter 21

Imagination, compass, and selection

Mini-roadmap. We now implement counterfactual search. IMAGINE may generate speculative nodes. COMPASS ranks them. SELECT chooses what to turn into an actual attempt. None of these functions can deposit a theorem directly.

Listing 21.1: Graded imagination

```
function IMAGINE(state, depth_budget, branch_budget):
    frontier = choose_future_frontier(state.futurebank)
    created = 0

    while not frontier.is_empty() and created < branch_budget:
        node = frontier.pop()
        if node.depth >= depth_budget:
            continue

        # Each generated action carries the role that generated it.
        actions = GENERATE_DEDUCTIVE_ACTIONS(state, node)

        for action in actions:
            if created >= branch_budget:
                break

            child = APPLY_COUNTERFACTUALLY(node, action)
            local_result = LOCAL_CHECK(action, node, child)
            if local_result == PASS:
                child.trust_grade = LOCAL
            elif action.is_allowed_as_explicit_hypothesis():
                child.trust_grade = SPEC
            else:
                continue # malformed or disallowed transitions are not branches

            child.id = fresh_id()
            child.parent_id = node.id
            child.generating_action = action
            child.role = action.role
            child.depth = node.depth + 1
            child.bank_dependencies = trace_bank_dependencies(action)
            child.assumed_dependencies = trace_hypotheses(action)
            child.estimated_cost = estimate_cost(state, node, action)
            state.futurebank.add(child)
            frontier.push(child)
```

```

        created += 1

    return state.futurebank

```

Listing 21.2: Compass and controller

```

function COMPASS(state):
    # K1 scores transitions, so the root (which has no incoming transition)
    # is not assigned a K1 edge score here.
    for node in state.futurebank.rankable_nonroot_nodes():
        parent = state.futurebank.parent_proof_state(node) # unique history parent
        node.compass_score = K1(parent, node.proof_state)

    actions = controller_actions(state)
    return argmax(actions, key=lambda a: K2(state, a))

function SELECT(state, controller_action, fallback=None):
    if controller_action.kind == "TRY_FUTURE_NODE":
        node = controller_action.node
        node.attempted = true
        proposal = materialize_certificate_proposal(node)
        proposal.source_future_node_id = node.id
        return proposal
    if controller_action.kind == "BACKFILL":
        return make_subgoal_proposal(controller_action.missing_dependency)
    if controller_action.kind == "SWITCH_ROLE":
        return next_agent_proposal(controller_action.role, state)
    if controller_action.kind == "FALLBACK" and fallback is not None:
        return fallback.next_proposal(state.bank)
    return NONE

```

21.1 Controlled structural escape

A controller can maintain a stagnation statistic based on repeated verifier failures, rank disagreement, or lack of bank growth. When tension is high it may raise novelty or switch regions. The control law is a search policy, not a theorem rule.

Chapter 22

The full AMLD loop

Mini-roadmap. This is the entire architecture in one place. The loop makes visible where each concept belongs and, just as importantly, where it does not belong.

Listing 22.1: Full AMLD loop

```
function AMLD_DECIDE(env, resource_budget):
    B0 = SEED_BANK(env)
    x0 = initial_proof_state(env, B0)
    F0 = new FutureBank(root = FutureNode(
        id = fresh_id(),
        proof_state = x0,
        parent_id = NONE,
        generating_action = NONE,
        role = NONE,
        assumed_dependencies = {},
        bank_dependencies = ids(B0),
        depth = 0,
        trust_grade = LOCAL
    ))

    state = SearchState(
        bank = B0,
        futurebank = F0,
        history = [],
        reflection_state = initialize_reflection(env),
        controller_state = initialize_controller(),
        step = 0
    )

    fallback = initialize_complete_fallback_if_available(env)

    while resource_budget.allows_next_step():
        state.step += 1

        # 1. Counterfactual search may grow without changing theorem memory.
        state.futurebank = IMAGINE(
            state,
            depth_budget = controller_depth(state),
            branch_budget = controller_width(state)
        )
```



```

# 2. Reflection may update correct self-descriptions used by control.
state.reflection_state = REFLECT(state)

# 3. Fair scheduling can reserve recurring fallback slots.
if fallback is not NONE and \
    is_reserved_fallback_step(state.step, state.controller_state):
    # Required to return the same typed Proposal used by all search paths.
    proposal = fallback.next_proposal(state.bank)
else:
    action = COMPASS(state)
    proposal = SELECT(state, action, fallback=fallback)

if proposal is NONE:
    continue

# 4. The verifier receives no compass score and no creativity score.
result = VERIFY(env, state.bank, proposal)

# 5. Only verifier-gated update can enlarge verified mathematical memory.
state, accepted_record = UPDATE(state, proposal, result)

# 6. Classify the certificate that was just accepted. Do not infer
# terminal semantics from whichever retrieval/index entry changed last.
if accepted_record is not NONE:
    status = TERMINAL_STATUS(env, state.bank, accepted_record)
    if status != NONE:
        return status, state

# 7. Search learns from success and failure without rewriting proof semantics.
state.controller_state = LEARN_CONTROL(state)

return UNKNOWN, state

```

A reader should be able to point to every safety boundary in this program. IMAGINE can be wrong. COMPASS can be wrong. REFLECT can be incomplete. SELECT can be foolish. The controller can waste time. Those are search failures. The logical contract says that none of them may directly add a verified theorem.

A fallback engine is a search component, not a fifth semantic role. Its `next_proposal` interface is required to return the same typed `Proposal` record as every other search path, including one submitting-role tag, a declared semantic certificate kind, an explicit claim and scope, and the certificate body. Fallback therefore has no separate route around `VERIFY`.

Chapter 23

Safe compression and quotienting

Mini-roadmap. FUTUREBANKs grow rapidly. We therefore want to merge states. The dangerous shortcut is to merge two histories merely because their visible endpoint formulas look the same.

Let $Q(h)$ be a compressed descriptor of FUTUREBANK history h .

Definition 23.1 (Future-sufficient descriptor). A descriptor Q is *future-sufficient* for a declared task if

$$Q(h_1) = Q(h_2)$$

implies equality of all continuation information relevant to the task's future search, verification, scope, and required counting/provenance semantics.

Proposition 23.2 (Safe FUTUREBANK quotient principle). *If Q is future-sufficient and any task-relevant path multiplicity or ancestry is retained separately, then histories with equal Q -state may share downstream computation without changing the represented continuation structure relevant to that task.*

Proof. By future-sufficiency, equality of descriptors implies equality of all continuation information that the declared task uses for future search, verification, scope, and counting/provenance semantics. Sharing the downstream computation therefore identifies only continuations that the task already treats as equivalent. Any path information that the task still distinguishes is preserved separately by hypothesis. $\square$

The difficult part is proving future-sufficiency.

Example 23.3 (Why endpoint equality fails). Two branches may display the same current goal g while one branch depends on a temporary assumption h and the other does not. Their visible endpoint formula is identical, but their legal future deposits differ. Merging them while discarding scope is unsound.

Chess supplies the analogous warning: two identical piece arrangements can differ in castling rights or repetition history, so future legality can depend on the path.

Design principle. Preserve provenance until you have proved which provenance is future-irrelevant.

Chapter 24

Software invariants and tests

Mini-roadmap. The theory suggests executable properties. A mature AMLD implementation should continuously test them, not merely describe them in prose.

Listing 24.1: Core invariant tests

```
property rejected_never_enters_bank:
  for every rejected proposal p:
    assert no r in bank.derived_records \
      where r.source_proposal_id == p.id

property every_derived_record_has_receipt:
  for r in bank.derived_records:
    assert r.verifier_receipt is not NONE
    assert r.source_proposal_id is not NONE

property verifier_policy_neutrality:
  choose same env, same bank, same semantic certificate s
  wrap s in proposals with different submitting roles
  vary futurebank, history, reflection, controller
  assert the intrinsic checker result on s is unchanged

property scope_is_dependency_closed:
  for r in bank.derived_records:
    assert all dependencies are admissible in r.scope

property futurebank_promotion_is_verifier_gated:
  for r in bank.derived_records:
    if r.source_future_node_id is not NONE:
      assert r.verifier_receipt is not NONE
      assert futurebank.promotion_of(r.source_future_node_id) == r.id

property quotient_preserves_required_provenance:
  for every merged quotient state q:
    assert task_required_ancestry(q) is reconstructible

property fair_fallback_progress:
  if fallback is enabled for an unbounded run:
    assert reserved schedule gives infinitely many effective slots
```

These properties are part of what it means to make the architecture scientific. They turn slogans

into falsifiable software claims. In particular, FUTUREBANK membership itself is never compared with BANK membership as though the two stored the same record type; promotion is tested through the explicit source-node/proposal/receipt chain.

Exercise 24.1 (Invariant violation). Suppose a developer modifies `VERIFY` to accept a certificate whenever either the formal checker accepts *or* the compass score exceeds 0.999. Identify every theorem or invariant in this text that the change threatens. Give the smallest architectural repair.

Part VII

Worked Examples and Exercises

Chapter 25

Four tiny worked AMLDs

Roadmap. We revisit each terminal role using bank states and certificates explicitly. These examples are deliberately elementary so that no logical subtlety hides the architecture.

25.1 P settles q from p and $p \rightarrow q$

Environment:

$$\Gamma = \{p, p \rightarrow q\}, \quad c = q.$$

Seed:

$$B_0 = \{p, p \rightarrow q\}.$$

Proposal: P supplies a modus-ponens certificate. Verification accepts. Deposit q . Terminal classifier recognizes $\Gamma \vdash c$ and returns PROVED.

The FUTUREBANK could have been empty. Imagination is an optimization layer, not a prerequisite for the logical meaning of this run.

25.2 R settles r from $\neg r$

Environment:

$$\Gamma = \{\neg r\}, \quad c = r.$$

The seed bank already contains the formula required by the refutation certificate. Depending on the certificate format, R may submit a trivial reference-to-premise proof. The accepted terminal status is REFUTED.

25.3 I settles p over empty propositional theory

Environment:

$$\Gamma = \emptyset, \quad c = p.$$

The I role supplies two valuations:

$$v_+(p) = \text{true}, \quad v_-(p) = \text{false}.$$

A small model checker verifies that v_+ satisfies $\Gamma \cup \{p\}$ and v_- satisfies $\Gamma \cup \{\neg p\}$. Using the declared soundness rule, the metacertificate concludes that neither side follows from Γ .

This example demonstrates why I is not simply “wait for P and R to fail.” It supplies a positive object.

25.4 C detects a broken foundation

Environment:

$$\Gamma = \{s, \neg s\}, \quad c = t.$$

The contradiction agent certifies $\perp$. In an explosive environment the system returns **INCONSISTENT** without treating later proofs of t as informative about the intended target status.

Chapter 26

A line-by-line PA AMLD run: $1 + 1 = 5$

Roadmap. This chapter lets one small false arithmetic equation pass through the whole architecture. We write the AMLD as D , fix $\Gamma = \text{PA}$, display the relevant seed BANK records, let all four roles act under a fair cyclic schedule, show what does and does not enter the FUTUREBANK and verifier history, and then inspect the winning R -certificate line by line. The example is intentionally elementary: the point is to make the machine visible.

26.1 The decider and its target

For this toy example, fix Peano arithmetic as the background theory and write

$$D(c) \in \{\text{PROVED}, \text{REFUTED}, \text{INDEPENDENT}, \text{INCONSISTENT}, \text{UNKNOWN}\}$$

for the operational status returned by the AMLD on a PA well-formed formula c . The letters P, R, I, C name search roles and designated settlement roles; they are not themselves status values. Thus a verified negated-target proof has designated settlement role R and licenses status **REFUTED**, regardless of which role submitted it. Under the contradiction-dominance convention, an accepted contradiction certificate licenses **INCONSISTENT** with priority in an explosive inconsistent background. **UNKNOWN** remains an operational resource outcome rather than a fifth logical truth class.

Let S denote successor. Then

$$1 = S0, \quad 2 = SS0, \quad 5 = SSSSS0.$$

Take the target

$$c: \quad S0 + S0 = SSSSS0.$$

After normalizing the addition on the left, this becomes the equation suggested informally at the start:

$$SS0 = SSSSS0.$$

The AMLD will settle the original formula by producing a checked proof of its negation.

26.2 The relevant slice of $\Gamma = \text{PA}$

The actual seed BANK contains the declared PA environment, not merely four formulas. For this proof, however, only a small slice is needed. Fix a standard presentation of PA (or a definitionally equivalent presentation) in which the following formulas are available as axioms or already verified derived laws, together with ordinary equality rules:

$$\begin{aligned} A_1 : \quad & x + 0 = x, \\ A_2 : \quad & x + S(y) = S(x + y), \\ A_3 : \quad & Sx \neq 0, \\ A_4 : \quad & Sx = Sy \rightarrow x = y. \end{aligned}$$

The environment also supplies universal instantiation/substitution, equality symmetry, transitivity and congruence, and classical logical rules. Induction is not needed for this particular target.

Thus

$$\mathbf{B}_0 = \text{seed}(\text{PA}, \mathcal{E}).$$

The displayed formulas are only the portion of $\mathbf{B}_0$ that the forthcoming certificate actually touches. Their origin is ENV, not one of the four agents.

26.3 A fair schedule chosen to make every role visible

A literal P, R, I, C round-robin would let R settle the target before the later roles necessarily do anything interesting. For pedagogy we instead use the cyclic order

$$P, I, C, R, P, I, C, R, \dots$$

which is just as fair. If no terminal certificate appeared, every role would receive infinitely many turns.

Initialize

$$\Sigma_0 = (\mathbf{B}_0, \mathbf{F}_0, \mathbf{H}_0, \text{Ref}_0, \text{Ctrl}_0), \quad \mathbf{F}_0 = \{h_0\}, \quad \mathbf{H}_0 = \emptyset,$$

where h_0 is the length-zero root history whose endpoint is the initial proof-search state induced by $\mathbf{B}_0$. No counterfactual branch has yet been added. We now follow one cycle.

26.4 Turn 1: P normalizes the left side

P is trying to prove c . It does not succeed, but it notices an ordinary arithmetic simplification that is useful to everyone.

Instantiate A_2 with $x = S0$ and $y = 0$:

$$S0 + S0 = S(S0 + 0).$$

Instantiate A_1 with $x = S0$:

$$S0 + 0 = S0.$$

By congruence of S ,

$$S(S0 + 0) = SS0.$$

By transitivity,

$$L_1 : \quad S0 + S0 = SS0.$$

P submits L_1 as

$$(P, (\text{INTERMEDIATE_THEOREM}, \kappa_1)).$$

The intrinsic verifier checks the arithmetic/equality derivation and returns **ACCEPT**. Therefore

$$B_1 = B_0 \cup \{L_1\}.$$

The record for L_1 has origin P , semantic kind **INTERMEDIATE_THEOREM**, scope Γ , and dependencies pointing back to the relevant PA seed records.

This is an important AMLD event: P has not proved the target, yet its work survives and may help an opposing role.

26.5 Turn 2: *I* considers independence

I must not infer independence from the absence of a proof. A natural speculative strategy is to seek a model of

$$PA \cup \{c\}$$

and a model of

$$PA \cup \{\neg c\}.$$

At this moment the role may place a plan of that form into the speculative **FUTUREBANK**:

$$f_I : \quad \text{“seek an admissible model-pair certificate for } c\text{.”}$$

No such terminal certificate has been verified, so the **BANK** does not change:

$$B_2 = B_1.$$

Verifier/search history records that I considered or attempted this route. The distinction is deliberate:

$$\text{interesting search information} \neq \text{proved mathematics}.$$

26.6 Turn 3: *C* contributes a theorem without finding inconsistency

C searches for a contradiction in the background. It does not produce a certificate of $PA \vdash \perp$. But while examining the successor axioms, it obtains a useful instance of A_3 by substituting $x = SS0$:

$$L_2 : \quad SSS0 \neq 0.$$

The verifier accepts the certificate for L_2 , so

$$B_3 = B_2 \cup \{L_2\}.$$

The origin field says C . The semantic kind is still an ordinary intermediate theorem. By Corollary 8.1, R may use it if its type and scope are admissible.

Again the role did not achieve its own terminal objective, yet its verified work remains useful.

26.7 The FUTUREBANK just before the winning turn

At this point a small imagined search structure might contain branches such as

- f_P : continue searching for a proof of c ,
- f_I : seek a model pair for independence,
- f_C : continue testing the PA background for contradiction,
- f_R : use L_1 to reduce the target to $SS0 = SSSSS0$ and strip successors.

A COMPASS is permitted to rank f_R highly because it now has verified dependencies L_1 and L_2 . That ranking does not itself prove anything. It merely helps choose the next attempt.

26.8 Turn 4: R submits the terminal refutation

R wants a certificate of

$$\neg(S0 + S0 = SSSSS0).$$

The following natural-deduction presentation makes the certificate readable line by line. A concrete PA implementation could encode the same reasoning in another proof calculus.

Line	Formula	Justification
1	$S0 + S0 = SSSSS0$	temporary assumption c
2	$S0 + S0 = SS0$	BANK lemma L_1
3	$SS0 = S0 + S0$	symmetry of line 2
4	$SS0 = SSSSS0$	transitivity, lines 3 and 1
5	$S0 = SSSSS0$	successor injectivity A_4 , line 4
6	$0 = SSS0$	successor injectivity A_4 , line 5
7	$SSS0 = 0$	symmetry of line 6
8	$SSS0 \neq 0$	BANK lemma L_2
9	$\perp$	lines 7 and 8
10	$\neg(S0 + S0 = SSSSS0)$	discharge line 1

Thus R submits a proposal whose semantic kind is `NEGATED_TARGET_PROOF`. Its dependency DAG contains a lemma originating with P , a lemma originating with C , and seed records originating with `ENV`. The verifier checks the certificate independently of those origins and accepts it.

The terminal BANK state is therefore

$$B_4 = B_3 \cup \{\neg(S0 + S0 = SSSSS0)\},$$

and the decider returns

$$D(c) = \text{REFUTED} \quad (\text{designated settlement role } R).$$

The run halts. I 's speculative model-pair branch never becomes theorem memory, and C never claimed that PA itself was inconsistent.

26.9 The BANK ledger for the whole run

The logical growth can be summarized compactly:

State	New verified content	Origin	Semantic kind
B_0	PA seed records, including the displayed relevant axioms	ENV	AXIOM/HYPOTHESIS
B_1	$L_1 : S0 + S0 = SS0$	P	INTERMEDIATE_THEOREM
B_2	no new verified record	—	—
B_3	$L_2 : SSS0 \neq 0$	C	INTERMEDIATE_THEOREM
B_4	$\neg(S0 + S0 = SSSSS0)$	R	NEGATED_TARGET_PROOF

The corresponding verifier/search history H additionally records the accepted receipts for L_1, L_2 , the terminal R -certificate, and I 's unsuccessful or merely speculative work. This is exactly why H is not the BANK.

26.10 What the example teaches

1. The target is settled by R , but R does not have to discover every lemma in its own certificate.
2. A P -discovered theorem can help refute the target; a C -discovered theorem can do the same.
3. The winning semantic role is determined by what the accepted certificate proves, not by the provenance of every dependency.
4. FUTUREBANK branches may guide attention without entering theorem memory.
5. A tiny false arithmetic equation already exercises Γ , seed records, shared BANK, provenance, agents, imagination, COMPASS selection, verification, and terminal classification.

Corollary 26.1 (Cross-role terminal assembly). *Under origin-neutral admissibility, an accepted terminal certificate may depend on verified BANK records originating from the environment and from any mixture of the roles P, R, I, C . If the dependencies are semantically admissible and the terminal certificate verifies, the terminal settlement meaning is determined by the certificate's semantic kind, not by the origin labels on its dependency DAG.*

Proof. Corollary 8.1 makes admissible verified records reusable across roles, while Corollary 3.8 makes terminal semantics independent of the discovering role. Combining the two facts gives the result. $\square$

Lemma 26.2 (Normalization-refutation lemma). *In classical first-order logic with equality, if*

$$\Gamma \vdash t = u \quad \text{and} \quad \Gamma \vdash u \neq v,$$

then

$$\Gamma \vdash t \neq v.$$

Proof. Assume $t = v$. From $t = u$, symmetry gives $u = t$, and transitivity with $t = v$ yields $u = v$, contradicting $u \neq v$. Discharge the temporary assumption. $\square$

The PA run is an instance with

$$t = S0 + S0, \quad u = SS0, \quad v = SSSSS0.$$

The BANK lemma L_1 supplies the normalization $t = u$; the rest of the R -certificate supplies $u \neq v$.

Exercise 26.3 (Replay the PA refutation). Repeat the run with target

$$c' : \quad SS0 = SSSSS0$$

so that no addition-normalization lemma is needed. Give the shortest readable R -certificate you can using successor injectivity and the axiom that no successor equals zero. Then explain which architectural features disappear from the example when L_1 is no longer needed.

Chapter 27

A worked FUTUREBANK example

Mini-roadmap. We now add imagination to a toy proof. The example shows the difference between seeing a route and owning the lemmas required to travel it.

Suppose Γ contains rules sufficient for

$$\lambda \rightarrow \mu, \quad \mu \rightarrow c,$$

but λ has not yet been proved. A speculative FUTUREBANK may contain the conditional route

$$[\text{suppose } \lambda] \rightarrow \mu \rightarrow c.$$

This is useful because it identifies λ as a strategically valuable subgoal. Yet neither μ nor c may be deposited globally from this branch while λ remains hypothetical.

The controller can backfill:

1. search for a certificate of λ from Γ ;
2. if accepted, deposit λ ;
3. apply the verified rule to obtain μ and deposit it after checking;
4. derive c and terminate after checking.

The imagined route functions like a chess variation: it tells the machine what a successful continuation would look like. Verification turns the variation into a realized proof history.

Chapter 28

Exercises

Mini-roadmap. These exercises are not research conjectures. They test whether the architecture is understood. Solution sketches follow in the next chapter.

Exercise 28.1 (Axiom versus theorem). Why should a seed record for $\gamma \in \Gamma$ not be labeled simply “proved”? Give a data schema that distinguishes an admitted premise from a derived theorem while allowing both to be used as premises when appropriate.

Exercise 28.2 (Independence certificate). In classical propositional logic with $\Gamma = \{p \vee q\}$ and target p , give two valuations that witness that p is independent of Γ . Explain what the verifier must check.

Exercise 28.3 (FUTUREBANK contamination). A speculative node contains the statement L because the controller asked “what if L were true?” A programmer inserts L into the BANK to make downstream search faster and adds a flag `speculative=true`. Is this safe if ordinary deduction routines can still read L ? Explain.

Exercise 28.4 (Fairness). Suppose complete fallback is invoked on steps 1, 2, 4, 8, 16, ... Is the schedule fair to fallback in the sense required by Theorem 16.1? Does the density of fallback steps need to stay bounded away from zero?

Exercise 28.5 (Unsafe quotient). Construct two proof histories that end with the same visible goal but should not be merged because one carries an extra local assumption. What extra information would make a quotient descriptor safer?

Exercise 28.6 (Role cooperation). Describe a run in which I discovers a model or structural fact that leads P to a useful object-level lemma, but I never produces a final independence certificate. What should be stored in BANK and what should remain in verifier/search history?

Chapter 29

Solution sketches

Mini-roadmap. The sketches emphasize the conceptual point of each exercise rather than presenting one mandatory implementation.

Exercise 3.5. Both proposals have submitting-role tag I because I discovered and submitted them. The ordinary lemma should carry an object-level intermediate semantic kind such as `INTERMEDIATE_THEOREM`; if accepted, it enters theorem memory with origin I and its ordinary logical scope/dependencies. The later proof of c carries semantic kind `TARGET_PROOF`. Its accepted BANK record still has origin I , but its terminal settlement role is P and it licenses `PROVED`. Provenance answers who found it; semantics answers what it proves.

Exercise 1.10. (a) computational stopping only; (b) both a computational terminal event and mathematical settlement if the checker is sound; (c) computational stopping only; (d) mathematical settlement and normally terminal computation; (e) neither mathematical settlement nor necessarily a justified exhaustive claim.

Exercise 2.2. B_0 contains admitted records for a , $a \rightarrow b$, and $b \rightarrow d$. The requested three BANK states are B_0, B_1, B_2 : start with admitted records for a , $a \rightarrow b$, and $b \rightarrow d$; verify and deposit b to obtain B_1 ; then verify and deposit d from b and $b \rightarrow d$ to obtain B_2 . Thus there are two derived deposits after the seed state. The three Γ records are admitted premises; b and d are derived.

Exercise 6.2. The certificate for u establishes at most $\Gamma \cup \{h\} \vdash u$. Re-labeling it as $\Gamma \vdash u$ silently discharges an assumption. Store an explicit scope or assumption set in each record and require the verifier to check dependency closure.

Exercise 8.2. Many answers work. For example, while attempting $\neg c$, R may prove a general algebraic identity L from Γ . If L is globally scoped and verified, P may use it later. Origin is provenance, not ownership.

Exercise 9.1. Position $\leftrightarrow$ proof state; legal move $\leftrightarrow$ locally admissible deduction; principal variation $\leftrightarrow$ favored proof route; transposition $\leftrightarrow$ two histories reaching a future-equivalent proof

state; evaluation $\leftrightarrow$ compass score; played move $\leftrightarrow$ selected and committed verified action. The analogy is imperfect because ATP branches may intentionally contain unproved hypothetical lemmas.

Exercise 12.2. Use a DAG $\lambda \rightarrow \mu \rightarrow c$. Speculative exploration may reason conditionally from both unproved nodes. Promotion requires an accepted certificate of λ , followed by an accepted certificate of μ using the now-verified λ , followed by the target.

Exercise 15.3. Examples: a simple nonreflective chess-like proof search (high FUTUREBANK, low formal self-description); a tagged trivial loop proving one self-description (low FUTUREBANK, high self-description); an ordinary simple prover (low/low); and a reflective FUTUREBANK controller (high/high). This shows that the capabilities are not one scalar.

Exercise 24.1. The change violates policy-neutral verification, noncontamination, and the principle that compass confidence cannot authorize theorem status. Repair: remove compass score from the acceptance disjunction and keep it only in selection/ranking.

Exercise 28.1. Use a field such as `kind=AXIOM|HYPOTHESIS|DERIVED`. Seed records obtain their authority from the environment declaration, while derived records require verifier receipts and dependencies.

Exercise 28.2. For $p \vee q$, choose $v_1(p) = T, q = F$ and $v_2(p) = F, q = T$. Both satisfy Γ ; one satisfies p , the other $\neg p$. The verifier checks the valuations and the metatheoretic soundness rule that converts the model pair into nonderivability on both sides.

Exercise 28.3. It is unsafe if ordinary deduction can consume the flagged record as a premise. A flag is not quarantine unless admissibility routines enforce it. Keep the item in F_{spec} or require a type system that makes speculative records unusable by verified deduction.

Exercise 28.4. Yes. The powers-of-two schedule gives infinitely many fallback steps, so it is fair in the theorem's weak sense. Its asymptotic density is zero; positive density is not required by Theorem 16.1, though performance may be poor.

Exercise 28.5. One history may end at goal g under scope Γ and another at the same displayed g under $\Gamma \cup \{h\}$. A safer descriptor includes active assumptions, admissible dependency set, role/metatheory context, and other continuation-relevant state.

Exercise 28.6. Store any independently verified, correctly typed structural fact that is admissible for later use. Store unsuccessful independence attempts, heuristic scores, and rejected metacertificates in history rather than theorem memory.

Exercise 18.12. The sorted probabilities are already $(0.40, 0.25, 0.15, 0.10, 0.10)$, so $G_1 = 0.40$, $G_2 = 0.65$, and $G_3 = 0.80$. With no new side information, any one-, two-, or three-candidate strategy can cover at most that much prior probability mass. More internal computation can reorder candidates, but it cannot make a three-element guessed set contain more than the top-three mass.

Exercise 18.13. The server-side retry counter is consumed only by interactions with the live server. A separate local verifier need not consult that state, so the number ten has no direct force there. The relevant defenses become the password distribution and the cost of each local verification attempt, which is why modern password hashing, unique salts, and suitable work factors matter.

Exercise 18.14. One suitable one-sided safety target is $c : \forall z[\text{Impl}(z) \Rightarrow (\text{Pwd}(z) \wedge \text{Factor}(z))]$. The refuter R needs a checked event z^* for which the implementation accepts but at least one required predicate is false. Such a witness is a finite counterexample certificate. A separate reverse implication would be needed if the exercise also required complete acceptance of all intended events.

Exercise 18.15. For example, a system that permanently locks an account after one failed attempt gives an unauthorized guesser almost no online budget, but any person who can submit one incorrect attempt for a known username may also deny the legitimate user access. The objective therefore includes both confidentiality/authentication risk and availability. Smaller N improves one dimension only under additional assumptions.

Exercise 26.3. Starting from $SS0 = SSSSS0$, successor injectivity gives $S0 = SSSS0$, then $0 = SSS0$. Symmetry gives $SSS0 = 0$, contradicting the PA axiom instance $SSS0 \neq 0$. Discharging the initial assumption yields $\neg(SS0 = SSSSS0)$. Because the target is already normalized, the example no longer needs the P -origin lemma L_1 or the normalization-refutation step; it therefore shows less cross-role BANK cooperation than the $1 + 1 = 5$ version.

Exercise 17.16. (a) A finite subset mentions only finitely many inequalities $H > c_{n_i}$, so the standard model satisfies that subset when H is interpreted by any ordinary m larger than all listed n_i . (b) Compactness gives a model of the whole theory. (c) A model of $\text{ElDiag}(\mathbb{N})$ yields an elementary embedding of the named standard structure; the added inequalities force H to lie above every named standard natural and hence outside that embedded copy. (d) “standard” refers to the externally designated embedded copy of $\mathbb{N}$; the construction expresses unlimitedness by the external scheme $H > c_n$ for each standard n , without adding a special “standardness” predicate to PA. (e) Model existence is not an effective presentation of the model or a finite runtime bound.

Exercise 17.17. (a) $B_H^M(\varphi)$ uses only the arithmetic elementary-extension framework. (b) “Internal hyperfinite interval” belongs to the richer full-NSA framework. (c) Elementarity of the arithmetic extension is enough. (d) A literal internal hyperfinite BANK history again requires a full nonstandard enlargement carrying the relevant state objects, not merely a nonstandard model of arithmetic.

Exercise 17.18. By the second incompleteness theorem, consistent PA does not prove its own usual consistency sentence. Therefore $\text{PA} + \neg \text{Con}(\text{PA})$ is consistent and has a model A . In

$A, \neg \text{Con}(\text{PA})$ says that some p satisfies the arithmetized PA proof predicate for contradiction. Externally no standard p can do so, hence the witness is nonstandard. The failed step is the appeal from $A \models \exists p \text{Prf}_{\text{PA}}(p, \ulcorner \perp \urcorner)$ back to the standard structure $\mathbb{N}$: $\mathbb{N}$ is not assumed elementary in this arbitrary model. Unlimitedness of the witness does not repair that failure.

Exercise 17.19. (a) M contains the embedded standard naturals as well as nonstandard elements, so membership $H \in M$ alone says nothing about whether H dominates every standard natural. (b) Any ordinary PA proof has a standard finite code n , and unlimited H satisfies $n < H$, so that same code witnesses $B_H^M(\varphi)$. (c) No standard N dominates all standard naturals; in particular $N + 1 > N$. Theorem 17.12 sharpens this by showing that some standard PA theorem has no proof code below any prescribed finite cutoff. (d) A positive bounded proof witness implies by elementarity that an ordinary finite proof exists. In the I case, failure of both bounded proof-existence conditions yields an external nonderivability classification, not automatically a finite independence certificate. (e) D_H^M is a metatheoretically defined classification using a nonstandard bound, not a standard algorithm with a standard finite runtime.

Part VIII

Research Frontier

Chapter 30

Conjectures

Roadmap. Only now, after the architecture is defined and exercised, do we state research conjectures. Each is deliberately separated from the proved invariants above. The following are hypotheses about search efficiency, not changes to proof semantics.

Conjecture 30.1 (Reflective navigation advantage). For some nontrivial families of proof-search problems, a FUTUREBANK-reflective controller with correct access to selected facts about its own search state can reduce expected search cost relative to an otherwise matched FUTUREBANK/compass architecture lacking those reflective features, while leaving the intrinsic verifier unchanged.

Conjecture 30.2 (Graded imagination advantage). For some hard proof-search distributions, allowing quarantined speculative FUTUREBANK expansion before expensive final verification can reduce expected total search cost relative to an architecture that fully verifies every imagined edge before deeper expansion, while preserving verifier-gated noncontamination.

Conjecture 30.3 (Tension-triggered structural escape). On some problem distributions with deceptive local guidance, a controller that increases structural novelty in response to sustained disagreement between compass ranking and verifier outcomes can outperform a matched controller with a fixed novelty schedule.

Conjecture 30.4 (Backfill advantage). On some theorem families with expensive auxiliary lemmas, conditional exploration followed by selective dependency backfilling can reduce expected work relative to proving every candidate auxiliary lemma before exploring its downstream consequences.

Conjecture 30.5 (Moderate reflection optimum). For at least some finite-resource AMLD architectures, an intermediate degree of formal self-description can yield better search efficiency than either no reflection or substantially more expensive reflection, because useful controller information is obtained before metareasoning overhead dominates.

The last conjecture is intentionally weaker than the slogan “Level 2 is always optimal.” No theorem in this book establishes a universal creativity optimum at Level 2. Any such universal claim would be easy to attack by changing the cost model or constructing a problem on which reflection is irrelevant.

30.1 Password-domain research conjectures

The proved results in the password special case were deliberately information-theoretic or logical. The following stronger claims concern search efficiency and security-control design and should be treated as conjectures until tested under preregistered models.

Conjecture 30.6 (Audit-over-secret advantage). On some realistic classes of well-rate-limited authentication systems with low top- N exposure, allocating AMLD-style imagination and verification primarily to specification-conformance, recovery-route, and control-policy auditing yields greater defensive value per unit computation than allocating the same resources primarily to direct password-candidate ranking.

Conjecture 30.7 (Cross-route shared-BANK advantage). For authentication systems with multiple sufficient routes such as password login, recovery, remembered sessions, and second-factor flows, a provenance-preserving shared BANK of verified invariants and counterexample lemmas can reduce expected audit cost relative to matched isolated audits of each route, particularly for defects that arise only from route interaction.

Conjecture 30.8 (Adaptive flood-control frontier). For some threat and usability distributions, an adaptive flood-control policy using verified account and risk state can achieve a strictly better tradeoff between unauthorized-acceptance probability and legitimate-user denial than every fixed hard-lockout policy in a specified comparison class.

Conjecture 30.9 (Verifier-history repair advantage). For some authentication-audit distributions, retaining rejected defect hypotheses and failed repair attempts in verifier history, while keeping them out of the verified BANK, can improve subsequent repair selection without increasing false certification of defects or fixes.

30.2 Tempting stronger claims that should already be rejected

The following are not retained as conjectures because simple counterexamples undermine them:

- **Deeper FUTUREBANKs can never hurt.** False under finite memory or time: extra branches can consume the budget before a useful shallow branch is attempted.
- **More speculation can only help.** False when speculation overwhelms ranking, memory, or backfill resources.
- **Same endpoint means safe merge.** False when scope or other path-dependent continuation information differs.
- **A sufficiently good compass makes verification unnecessary.** False as an architectural claim: ranking confidence and certificate validity are different predicates.
- **Long joint failure of P and R establishes independence.** False: search failure is not a nonderivability certificate.
- **A certificate is valid because it came from the “right” agent.** False: under the narrow interface, agent identity is provenance; validity is determined by semantic certificate content, scope, dependencies, environment, and checker.

- **Higher formal self-awareness necessarily improves imagination.** False by Proposition 14.1; the capabilities are logically distinct.

Chapter 31

LAB: How to try to refute the conjectures

Roadmap. This Lab is adversarial by design. The reader’s job is not to make the conjectures look good. It is to give them a fair chance to fail. We first explain the general experimental protocol, then give a refutation roadmap for every conjecture above.

31.1 First lesson: a vague existential conjecture is hard to refute experimentally

Statements such as “there exist some problem families on which method A helps” are mathematically weak and scientifically slippery. Failure on one benchmark does not refute the existence of another benchmark where A helps.

Therefore the Lab converts each broad research conjecture into a preregistered benchmark-level hypothesis. Fix:

- a benchmark distribution D ;
- a train/dev/test split chosen before the final comparison;
- identical verifier V and verification environment format;
- identical total CPU/GPU/time or verifier-call budget;
- fixed random seed protocol;
- a primary metric, such as solved fraction, verifier calls to settlement, wall-clock time, or area under a solved-versus-budget curve;
- a minimum practically meaningful effect size δ ;
- a statistical decision rule that distinguishes “insufficient evidence” from evidence that a prespecified effect is smaller than δ .

Then ask a finite question. For example:

On preregistered test distribution D , does reflective controller A improve median verifier calls by at least δ relative to matched controller B?

That claim can fail cleanly.

31.2 Universal Lab protocol

For every conjecture:

1. **Freeze correctness.** Use exactly the same verifier binary/version and certificate semantics in every condition.
2. **Change one architectural feature.** Do not compare two systems that differ in five uncontrolled ways.
3. **Pair runs.** Run both conditions on the same problems and, where stochasticity exists, matched random seeds.
4. **Record the whole trace.** Keep BANK growth, verifier calls, FUTUREBANK node counts, rejected proposals, role allocations, fallback steps, and terminal status.
5. **Use a resource curve.** A method that wins only because it received more computation has not established the intended advantage.
6. **Attempt counterexample construction.** In addition to random benchmarks, design adversarial instances specifically meant to exploit the proposed mechanism's overhead.
7. **Publish negative results.** A clean refutation is scientific progress.

31.3 Lab A: Reflective navigation advantage

Ablation. Compare the same FUTUREBANK, compass, verifier, and fallback with reflection enabled versus disabled. Reflection-enabled search may read certified self-state facts; reflection-disabled search receives matched non-self-referential control inputs where possible.

Refutation target. Preregister a benchmark-relative claim such as a minimum δ reduction in verifier calls together with a one-sided confidence-bound or equivalence-testing rule capable of excluding an improvement of that size. Refute that finite claim when the preregistered analysis excludes an improvement of size δ (or supports a practically meaningful effect in the wrong direction). Merely failing to obtain statistical significance is not, by itself, evidence that the claimed advantage is absent.

Adversarial cases. Use problems where the correct proof is nearly immediate and reflection overhead is pure cost; use problems where self-state statistics are intentionally uninformative; use a controller that can already infer equivalent information without formal reflection.

31.4 Lab B: Graded imagination advantage

Ablation. Condition G permits F_{spec} expansion; condition E requires each edge to be fully certified before deeper expansion. Keep branch generator and ranking model matched.

Refutation target. Use the preregistered decision rule to show that speculative expansion falls short of the declared efficiency threshold, or that memory/branch explosion increases total cost enough to erase the benefit.

Adversarial cases. Construct domains in which almost every speculative auxiliary lemma is false or unprovable, so that deep counterfactual exploration spends effort on dead branches.

31.5 Lab C: Tension-triggered structural escape

Ablation. Compare an adaptive novelty controller using compass/verifier disagreement with a fixed novelty schedule matched for average novelty budget.

Refutation target. On a preregistered deceptive-guidance benchmark, use the declared decision rule to show that the tension signal fails to achieve the stated improvement, or that it systematically triggers escapes just before the local search would have succeeded.

Adversarial cases. Create noisy verifier-failure patterns that are not evidence of local trapping. If tension confuses ordinary hard progress with failure, it may thrash between regions.

31.6 Lab D: Backfill advantage

Ablation. Condition B explores conditional downstream utility before proving auxiliary lemmas. Condition P requires proof-first certification of each candidate lemma. Use the same candidate generator.

Refutation target. Under the preregistered resource metric, show that speculative downstream analysis costs more than the proof effort it saves, or that it repeatedly selects high-utility-looking lemmas that turn out to be impossible to backfill.

Adversarial cases. Generate many seductive hypothetical lemmas that would trivialize the target if true but are false or extremely hard to prove.

31.7 Lab E: Moderate reflection optimum

Ablation. Implement at least three reflection budgets: none, moderate, and high. If using the Level terminology, implement the declared self-description schemas explicitly and charge their construction/checking cost rather than treating a level label as free metadata.

Refutation target. Refute a benchmark-specific “middle is best” claim under the preregistered decision rule by establishing monotone improvement with more reflection, monotone degradation, or equivalence within the declared practical margin. To attack the stronger slogan “Level 2 is universally optimal,” a single well-specified architecture and benchmark in which another declared level strictly wins is enough.

31.8 Lab F: Audit-over-secret advantage

Ablation. Fix an authorized synthetic authentication model with a hard or effective attempt budget N . Give condition S most discretionary computation to candidate ranking and condition A the same budget for formal conformance, recovery-route, and policy-defect search. Keep the underlying model and checker fixed.

Refutation target. Reject the benchmark-level advantage claim when the preregistered analysis excludes the declared value improvement for condition A—for example, because verified defect/invariant yield is equivalent within the practical margin or because overhead erases the claimed gain.

31.9 Lab G: Cross-route shared-BANK advantage

Ablation. Compare isolated route audits with an otherwise matched shared-BANK audit. The shared condition may reuse only verified, correctly scoped invariants and lemmas.

Refutation target. Use the preregistered decision rule to establish that sharing provides less than the declared practical audit benefit, or that bookkeeping and overly broad reusable facts create enough overhead to make isolation better.

31.10 Lab H: Adaptive flood-control frontier

Ablation. Compare preregistered fixed lockout policies against an adaptive throttling controller on the same simulated legitimate-user and adversarial-attempt distributions.

Refutation target. Plot unauthorized-acceptance risk against legitimate-user denial and latency. Refute the finite benchmark claim if the preregistered analysis establishes that the adaptive policy is Pareto-dominated by a fixed policy or falls short of the declared improvement margin.

31.11 Lab I: Verifier-history repair advantage

Ablation. Compare two repair-search systems that share the same verified BANK. One may consult structured history of rejected defect hypotheses and failed repairs; the other may not.

Refutation target. Reject the advantage claim when the preregistered analysis excludes the declared improvement in verified repair rate, or when history-driven search produces a practically meaningful increase in wasted budget on stale failure modes.

31.12 What counts as a successful refutation?

There are two very different standards:

1. A **mathematical universal claim** can be refuted by one valid counterexample satisfying its hypotheses.

2. An **empirical benchmark claim** is refuted according to its preregistered finite decision rule.

Do not pretend that a failed experiment disproves an existential mathematical statement over all possible problem families. Instead, write benchmark-level hypotheses strongly enough that the data are allowed to say “no.”

31.13 Minimal experiment pseudocode

Listing 31.1: Paired ablation harness for any Lab conjecture

```
function RUN_ABLATION(benchmark, system_A, system_B, budgets, seeds):
    results = []

    for problem in benchmark.test_problems:
        for seed in seeds:
            for budget in budgets:
                trace_A = run(system_A, problem, seed, budget)
                trace_B = run(system_B, problem, seed, budget)

                assert trace_A.verifier_version == trace_B.verifier_version
                assert trace_A.problem_environment == trace_B.problem_environment

                results.append({
                    "problem": problem.id,
                    "seed": seed,
                    "budget": budget,
                    "A": summarize(trace_A),
                    "B": summarize(trace_B)
                })

    return paired_analysis(results, preregistered_primary_metric)
```

The final scientific habit is perhaps the most important one in the book: a search idea should earn its place by surviving both the verifier and the experiment appropriate to it.

Closing roadmap

The architecture can now be remembered in nine sentences.

1. Γ is the admitted starting point, and the BANK begins by recording it explicitly.
2. P, R, I, C are four principal search roles. Every agent proposal has exactly one submitting-role tag, while certificate semantic kind is a separate coordinate.
3. A terminal certificate has a designated settlement role determined by its verified semantics; its submitting role is provenance and need not match that settlement role.
4. The BANK is shared, typed, provenance-preserving verified memory, so intermediate mathematics can cross agent boundaries safely.
5. The FUTUREBANK contains represented possible proof futures, including quarantined hypothetical lemmas, much as chess search contains moves that have not been played.
6. A deductive move changes a proof-search state; imagination recursively explores such moves, while planning and the compass decide which imagined route deserves real computation.
7. Formal self-description may improve control, but it is a separate axis from imagination and may never redefine the verifier.
8. AMLD is a general architecture whose applications may include mathematical settlement, software and protocol verification, engineering safety, formalized scientific reasoning, and many other certificate-bearing domains; authentication and passwords are only one worked special case showing how domain-specific information and external controls can bound search.
9. The verifier is the fixed gate between “this might work” and “this is now an accepted formal conclusion in the declared environment.”

References and live links

- [1] Max Tegmark, “Is ‘the theory of everything’ merely the ultimate ensemble theory?,” *Annals of Physics* 270 (1998), 1–51. ArXiv: <https://arxiv.org/abs/gr-qc/9704009>. DOI: <https://doi.org/10.1006/aphy.1998.5855>.
- [2] Claude E. Shannon, “Programming a Computer for Playing Chess,” *The London, Edinburgh, and Dublin Philosophical Magazine and Journal of Science*, Series 7, 41(314) (1950), 256–275. DOI: <https://doi.org/10.1080/14786445008521796>.
Computer History Museum copy: <https://www.computerhistory.org/chess/doc-431614f453dde/>.
- [3] Metamath Proof Explorer, Theorem `2p2e4`, “Two plus two equals four.” <https://us.metamath.org/mpeuni/2p2e4.html>.
Metamath Proof Explorer home page: <https://us.metamath.org/mpeuni/mmset.html>.
- [4] Geoff Sutcliffe and the TPTP project, “TPTP and TSTP Quick Guide.” <https://tptp.org/UserDocs/QuickGuide/>.
- [5] TPTP project, “The TPTP Language” and SZS status documentation. <https://tptp.org/UserDocs/TPTPLanguage/TPTPLanguage.shtml>.
- [6] Max Wisniewski, *Agent-based Blackboard Architecture for a Higher-Order Theorem Prover*, Master’s thesis, Freie Universität Berlin, 2014. Public PDF: [live PDF](#).
- [7] Jan Jakubův and Josef Urban, “ENIGMA: Efficient Learning-based Inference Guiding Machine,” 2017. ArXiv: <https://arxiv.org/abs/1701.06532>.
- [8] Brian Tenneson, *Automated Logical Deciders: Certificate-Indexed Halting, Multiagent Search, Provenance-Aware Shared Lemma Pools, and Machine-Learned Guidance*, research supplement, August 2026. Snapshot repository: <https://github.com/btenneson/Automatic-Logical-Deciding-Snapshot-August-2-2026/>.
- [9] Brian Tenneson, *Depths of a Simulation & Formalized Self-Awareness*, merged edition, July 2026. Author manuscript. Public repository landing page: <https://github.com/btenneson/Depths-of-a-Simulation>.
- [10] Brian Tenneson, *Fund the Proof Machines—Guard the Capabilities That Can Become Weapons*, Substack essay, 2026. The essay argues for broad investment in open, independently verifiable ATP research while treating operationally dangerous capabilities as a separate governance problem; its dual-use examples include mathematics, software verification, authentication, cryptographic protocols, and cyber-physical systems. Live link: <https://btenneson2301.substack.com/p/fund-the-proof-machinesguard-the>.

- [11] National Institute of Standards and Technology, *Digital Identity Guidelines: Authentication and Authenticator Management*, NIST SP 800-63B-4, July 2025. Live HTML: <https://pages.nist.gov/800-63-4/sp800-63b.html>. DOI: <https://doi.org/10.6028/NIST.SP.800-63B-4>.
- [12] OWASP Cheat Sheet Series, “Authentication Cheat Sheet.” Live guide: [OWASP Authentication Cheat Sheet](#).
- [13] OWASP Cheat Sheet Series, “Password Storage Cheat Sheet.” Live guide: [OWASP Password Storage Cheat Sheet](#).
- [14] Herbert B. Enderton, *A Mathematical Introduction to Logic*, 2nd ed., Academic Press, 2001. See Section 2.5 for first-order Soundness, Completeness, and Compactness, Section 2.8 for nonstandard analysis, and Section 3.7 for arithmetized provability and the Second Incompleteness Theorem. DOI: <https://doi.org/10.1016/C2009-0-22107-6>. Publisher page: [ScienceDirect](#).
- [15] David Marker, *Model Theory: An Introduction*, Graduate Texts in Mathematics 217, Springer, 2002. See Section 2.1 for Compactness and Section 2.3 (beginning on p. 44) for elementary embeddings, diagrams, and elementary extensions. DOI: <https://doi.org/10.1007/b98860>.
- [16] H. Jerome Keisler, *Foundations of Infinitesimal Calculus*, online monograph, 2007 (online edition updated June 2022). University of Wisconsin–Madison author page and live monograph link: <https://people.math.wisc.edu/hkeisler/foundations.html>.
- [17] Robert Goldblatt, *Lectures on the Hyperreals: An Introduction to Nonstandard Analysis*, Graduate Texts in Mathematics 188, Springer, 1998. See pp. 35–58 for transfer and infinite/infinitesimal hyperreals, pp. 125–154 for internal/external and hyperfinite sets, and pp. 157–189 for nonstandard frameworks and existence. DOI: <https://doi.org/10.1007/978-1-4612-0615-6>.
- [18] Rineke Verbrugge, “Provability Logic,” *The Stanford Encyclopedia of Philosophy*, substantive revision 2024. Live entry: <https://plato.stanford.edu/entries/logic-provability/>.

Editorial provenance note

The AMLD roles, shared BANK, provenance emphasis, FUTUREBANK, compass/control architecture, verifier separation, and formalized self-description program are developed from Brian Tenneson’s prior manuscripts and research discussions. Versions 1.01–1.03 introduced and progressively qualified the compactness-based nonstandard discussion. Version 1.04 performs a stabilization audit of the entire nonstandard-analysis chapter: it makes the arithmetic elementary-extension framework primary, quarantines literal internal/hyperfinite language to a separately identified full-NSA framework, replaces process-like claims about an externally observable completed search with a bounded truth condition, proves model-and-cutoff invariance on standard targets, qualifies standard-part language, adds Goldblatt as a dedicated NSA reference, and supplies solution sketches for the new NSA exercises. Version 1.05 makes the framework separation syntactic as well as verbal: the compactness-built arithmetic extension is denoted M , the star notation $^*\mathbb{N}$ is reserved for an optional full NSA enlargement, and the bounded condition/classifier are written B_H^M and D_H^M . This removes the last systematic ambiguity between model-theoretic nonstandard arithmetic and full internal/hyperfinite NSA. Version 1.06 makes explicit the strongest hidden hypothesis: the reverse implication from a nonstandard proof witness to an ordinary proof depends on $\mathbb{N} \preceq M$, not merely $M \models \text{PA}$. A Gödel-II counterexample shows why arbitrary nonstandard PA models are insufficient, and the text now states that the full elementary diagram is a non-effective semantic resource rather than a computational recipe. Version 1.07 makes explicit that the compactness-built extension satisfies $M \models \text{Th}(\mathbb{N})$, not merely PA, and repairs the last dropped ambient-model superscript in the bounded-condition notation. Version 1.08 defines the elementary-extension notation at first use, makes the ambient membership $H \in M$ explicit in the trichotomy theorem, and spells out the initial-segment fact used in the Gödel-II counterexample. Version 1.09 is the solution-sketch fixed-point pass: it removes the last residual $^*\mathbb{N}$ notation from the arithmetic-only answer to Exercise 17.19(a), replacing it by the ambient elementary extension M . The chapter’s fixed correctness claims are now stated entirely in terms of standard PA syntax, elementary extensions, unlimited bounds, and metatheoretic classification; no theorem depends on treating a nonstandard bound as an executable finite resource. The book was substantially reorganized, expanded, formalized, technically audited, and edited with AI assistance across the v1.0–v1.09 development. Version 1.49 is the first post-NSA stable-release audit: Chapter 17 and its NSA solution sketches are held fixed, while the remaining text is checked for mathematical, architectural, statistical, notation, pseudocode, cross-reference, citation, and presentation defects. Version 1.89 is the second stable audit, resolving remaining global notation collisions, defining the full AMLD state space explicitly, sharpening the finite-attempt authentication model, aligning the illustrative pseudocode with the mathematical interfaces, and tightening empirical-refutation language. Version 2.29 repairs the remaining role/status conflation in the worked PA run and corrects a residual quotation defect, while continuing to hold the nonstandard-analysis material fixed. Version 2.69 repairs the executable interface contract by making proposal claims and scopes explicit checker inputs, binding accepted

receipts to those checked fields, adding the FUTUREBANK node identifier used by the pseudocode, and requiring fallback to return the same typed proposal object as every other search path. Version 3.09 removes two residual overstatements: it distinguishes environment-admitted seed assumptions from verifier-gated derived deposits, states backfill as a sufficient rather than necessary promotion condition, and sharpens the benchmark-refutation testing language. Version 3.49 distinguishes the authentication chapter's one-sided unauthorized-acceptance safety property from full extensional implementation/specification conformance, preventing a rejecting-all implementation from being described as fully conformant. Version 3.89 resolves a FUTUREBANK representation mismatch: transition scores and generating actions are now attached to single-parent history nodes in the core pseudocode, while DAG-style computation sharing is deferred to the separately proved safe-quotient layer. Version 4.29 repairs FUTUREBANK initialization and root handling: the full AMLD loop now seeds an explicit root history node from the initial BANK/proof state, and first-order transition scoring excludes that parentless root. Version 4.69 completes the pseudocode identity and budgeting contract: BANK records and proposals receive stable IDs, newly imagined history nodes receive their ID/parent/action/role/depth metadata explicitly, and IMAGINE now visibly enforces both depth and branch limits rather than referring to an unconsumed abstract budget. Version 5.09 repairs provenance identity across namespaces: derived BANK records now retain their source proposal ID, FUTUREBANK-materialized proposals retain their source node ID, rejected-proposal invariants test that actual link, and accepted promotion is recorded against the exact source future node. Version 5.49 fixes the remaining record-identity ambiguity by separating immutable BANK ledger records from formula-level retrieval deduplication, so repeated proofs can share lookup structure without erasing accepted receipts or provenance. Version 5.89 repairs the FUTUREBANK promotion invariant: accepted BANK records now retain the source FUTUREBANK node ID, SELECT explicitly marks and links attempted future nodes, and the executable test follows that provenance chain instead of comparing unlike FutureNode and BankRecord objects. Version 6.29 removes remaining cross-chapter notation collisions in the authentication example, aligns the worked PA run with the nonempty root-history definition of FUTUREBANK, and makes the seed exercise's two-deposit/three-state request internally consistent. Version 6.30 is a document fixed-point pass: it adds the author credentials, contact information, copyright notice, explicit page-number footer, and running-head apparatus; audits citation keys against bibliography entries; verifies the principal live reference targets; makes the previously uncited program-context and Keisler references explicit in the text; and rechecks LaTeX compilation, cross-references, bibliography resolution, and rendered-page layout until a subsequent audit pass produces no further source changes. Version 6.40 adds a plain-language elementary description to the title page so that the core distinction between imaginative search and verified mathematical acceptance is visible before any technical terminology is assumed. Version 6.80 expands that title-page description to a longer elementary blurb and adds a separate elementary foreword designed to give a non-specialist reader a two-to-three-page conceptual bridge into the BANK, FUTUREBANK, four settlement roles, verifier boundary, UNKNOWN output, COMPASS-guided search, provenance, and formalized self-description before the technical foreword begins. Because the scope and prose have changed considerably from earlier preview drafts, attribution percentages stated in those earlier drafts should not be assumed to apply unchanged to this edition.